

# **Fast algorithms exploiting low-rank structure for graph clustering and integral equations**

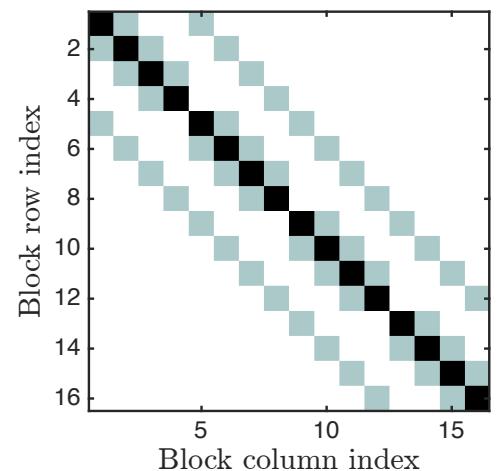
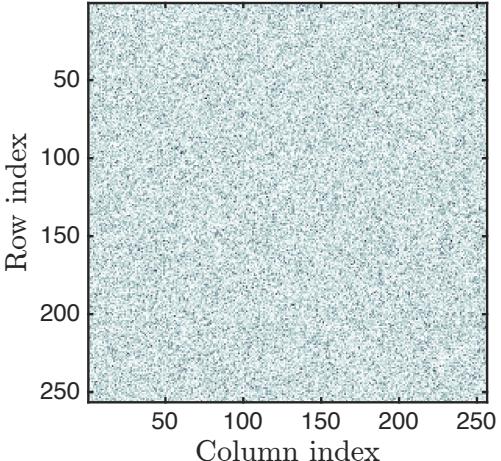
Victor Minden, Stanford ICME

Joint work with Anil Damle, Ken L. Ho, & Lexing Ying

December, 2016

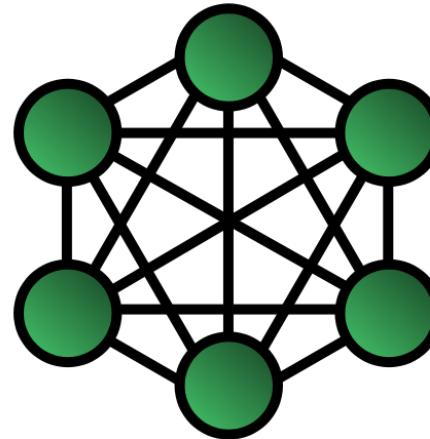
# Introduction

- My work: fast algorithms exploiting sparsity/related structure
- This talk
  - › Part 1: a simple linear algebraic algorithm for graph clustering
  - › Part 2: a rank-structured factorization for the inverse operator of fast multipole method systems (discretizations of integral equations)



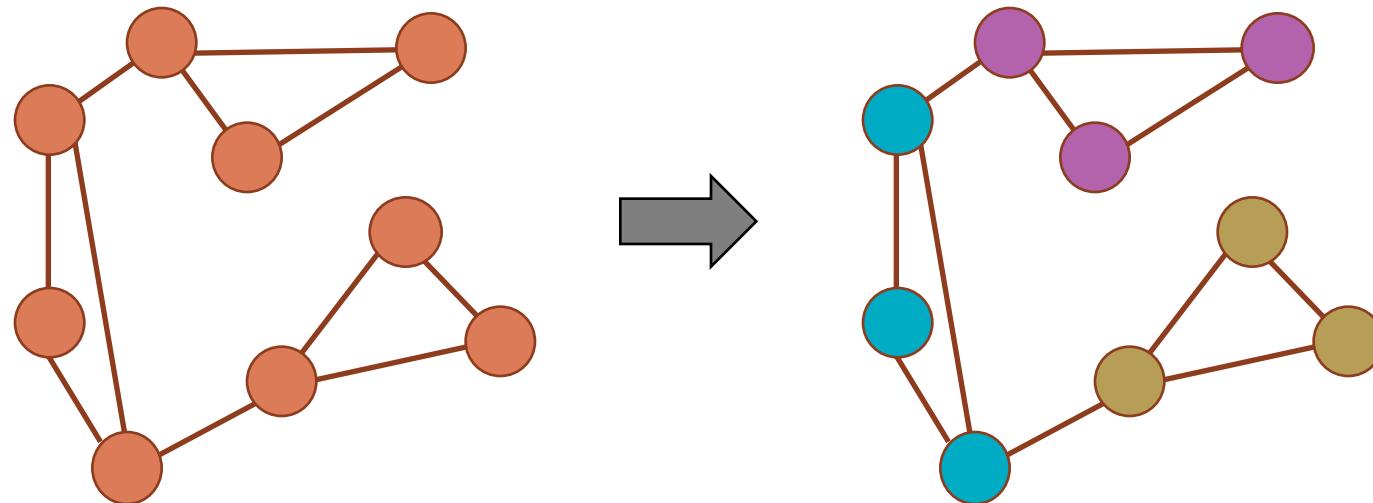
# **Robust and efficient multi-way spectral clustering**

**WITH DAMLE, YING**



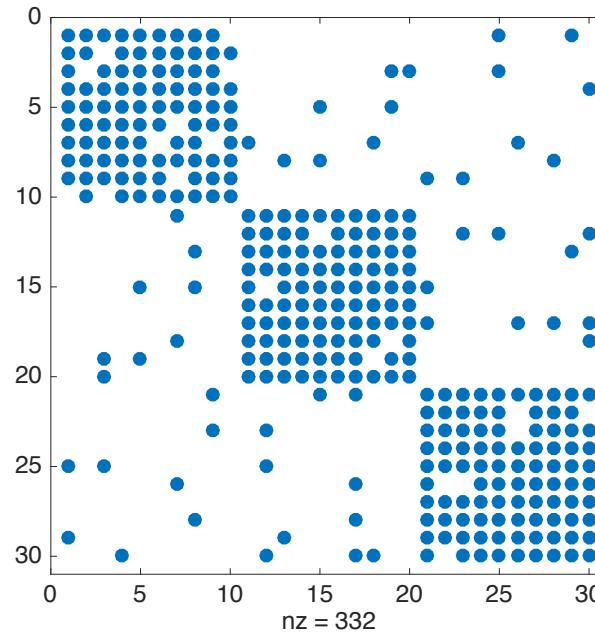
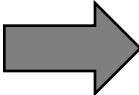
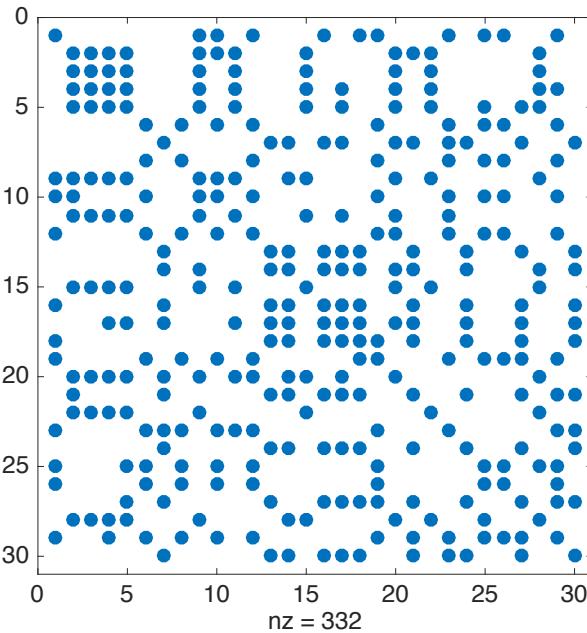
# Graph clustering

Task: find  $k$  relatively well-connected communities of nodes in a simple graph



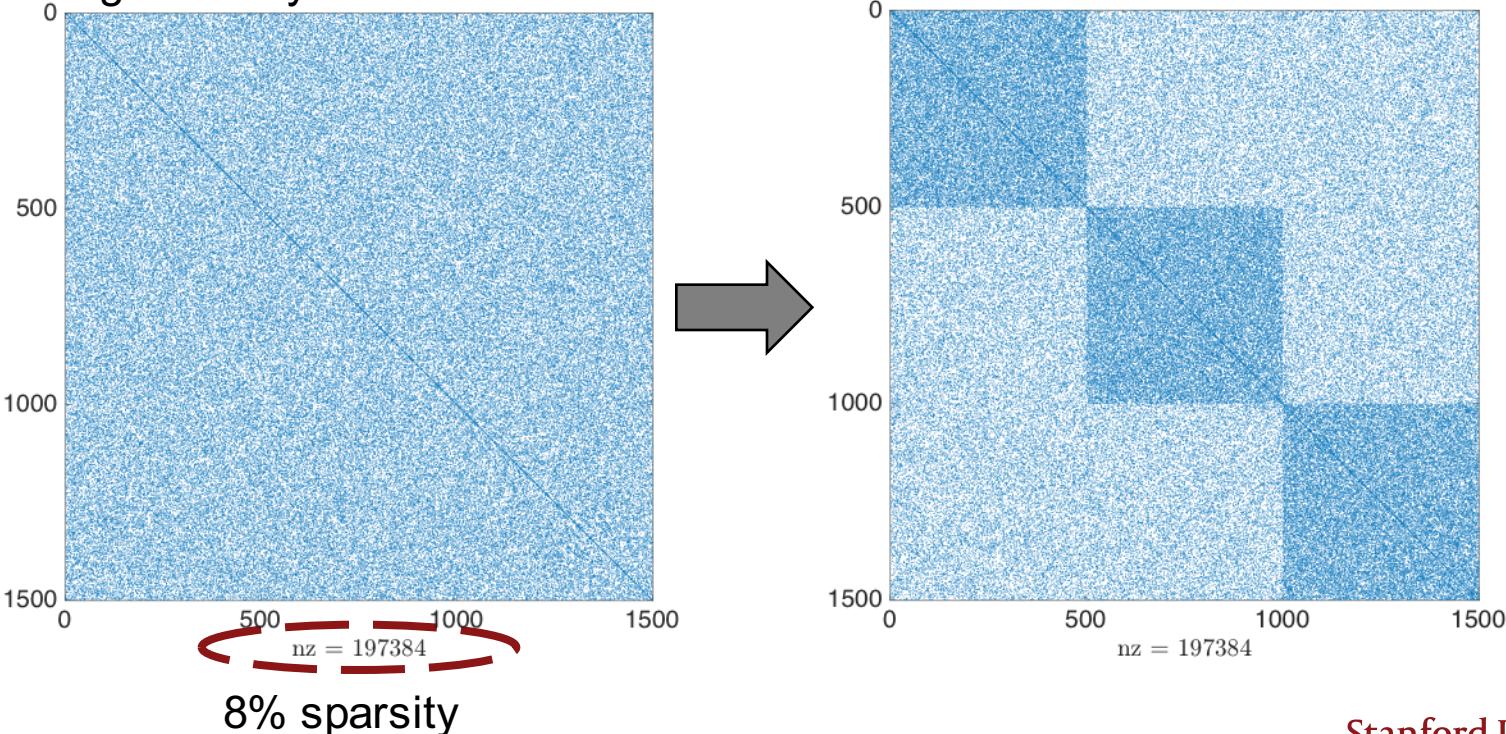
# Graph clustering

Sometimes this is easy, and adjacency matrix is essentially block diagonal under some permutation



# Graph clustering

Sometimes “within-cluster” edge density barely different from “between-cluster” edge density



## Basics

- For a graph  $G$  with  $n$  vertices, the (symmetric) adjacency matrix of  $G$  is  $A \in \{0,1\}^{n \times n}$  with

$$A_{ij} = \begin{cases} 1 & \text{edge } (i,j) \text{ is in } G \\ 0 & \text{else.} \end{cases}$$

- The stochastic block model (SBM) [Holland et al., 1983]
  - Entries of the symmetric adjacency matrix  $A$  are drawn independently from a Bernoulli distribution, where (for our purposes)

$$\mathbb{P}[A_{ij} = 1] = \begin{cases} p & i \text{ and } j \text{ in same cluster,} \\ q & \text{else.} \end{cases}$$

## The SBM mean adjacency matrix (k=3 clusters)

$$\mathbb{E}[A] = M = \Pi \begin{pmatrix} p & p & q & q & q & q \\ p & p & q & q & q & q \\ q & q & p & p & q & q \\ q & q & p & p & q & q \\ q & q & q & q & p & p \\ q & q & q & q & p & p \end{pmatrix} \Pi^T$$

$$= \Pi \begin{pmatrix} 1 & & & & & \\ 1 & 1 & & & & \\ & & 1 & & & \\ & & & 1 & & \\ & & & & 1 & \\ & & & & & 1 \end{pmatrix} \begin{pmatrix} p & q & q \\ q & p & q \\ q & q & p \end{pmatrix} \begin{pmatrix} 1 & & & & & \\ 1 & 1 & & & & \\ & & 1 & & & \\ & & & 1 & & \\ & & & & 1 & \\ & & & & & 1 \end{pmatrix}^\top \Pi^T = \Pi X \begin{pmatrix} p & q & q \\ q & p & q \\ q & q & p \end{pmatrix} X^T \Pi^T$$

## Relation to eigendecomposition of M

Using  $\begin{pmatrix} p & q & q \\ q & p & q \\ q & q & p \end{pmatrix} = Q\Lambda Q^T$ , we see eigenvectors are a **rotation** of indicators:

$$\mathbb{E}[A] = M = \Pi X \begin{pmatrix} p & q & q \\ q & p & q \\ q & q & p \end{pmatrix} X^T \Pi^T = (\Pi X Q) \Lambda (Q^T X^T \Pi^T) = U_k \Lambda' U_k^T$$

---

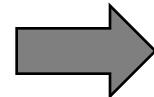
$$U_k = \Pi X Q$$

$\Pi X \in \mathbb{R}^{n \times k}$  is a collection of indicator vectors  
 $Q \in \mathbb{R}^{k \times k}$  is a rotation

## Another way to think about it

- Relax the maximum likelihood estimator for clusters

$$\begin{aligned} & \text{maximize} && \text{Trace}(X^T A X) \\ & \text{subject to} && X^T X = \frac{n}{k} I_k \\ & && X_{ij} \text{ binary} \end{aligned}$$



$$\begin{aligned} & \text{maximize} && \text{Trace}(X^T A X) \\ & \text{subject to} && X^T X = \frac{n}{k} I_k \\ & && X_{ij} \text{ real-valued} \end{aligned}$$

- We hope the solution to the relaxed problem is close to the solution to the binary problem.
- But, relaxed solution is determined only *up to rotation*:  $X^* = V_k Q$

## Our clustering algorithm

$$[Q, R, P] = qr(V^T k, 0);$$
$$[\tilde{c}, c] = \max(\text{abs}(R^* P'));$$

- Given matrix of top-k eigenvectors  $V_k$ , basic form has two steps:
  - Compute a column-pivoted QR factorization of  $V_k^T$

$$V_k^T \Pi = QR$$

- For each  $j = 1, \dots, n$ , assign vertex  $j$  to cluster

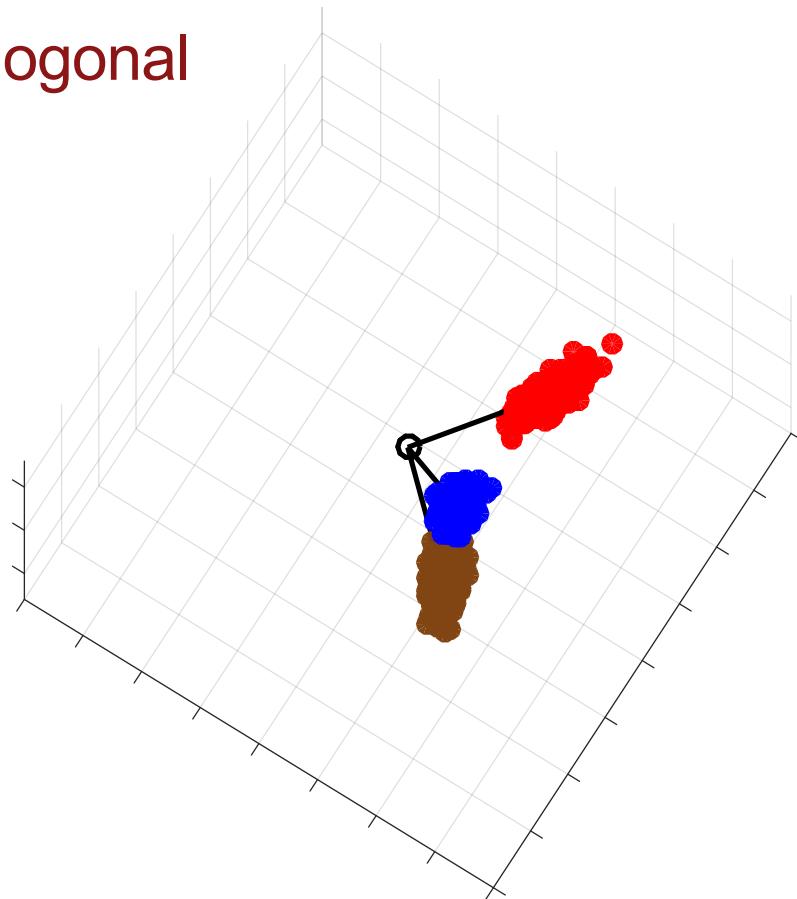
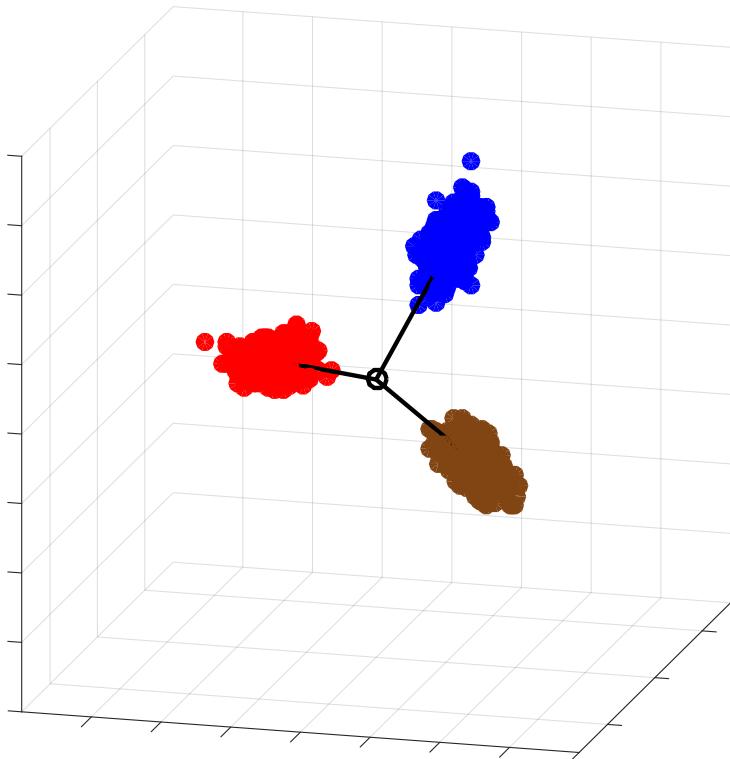
$$c_j = \arg \max_i |R\Pi^T|_{i,j}$$

---

- Complexity:  $\mathcal{O}(nk^2)$

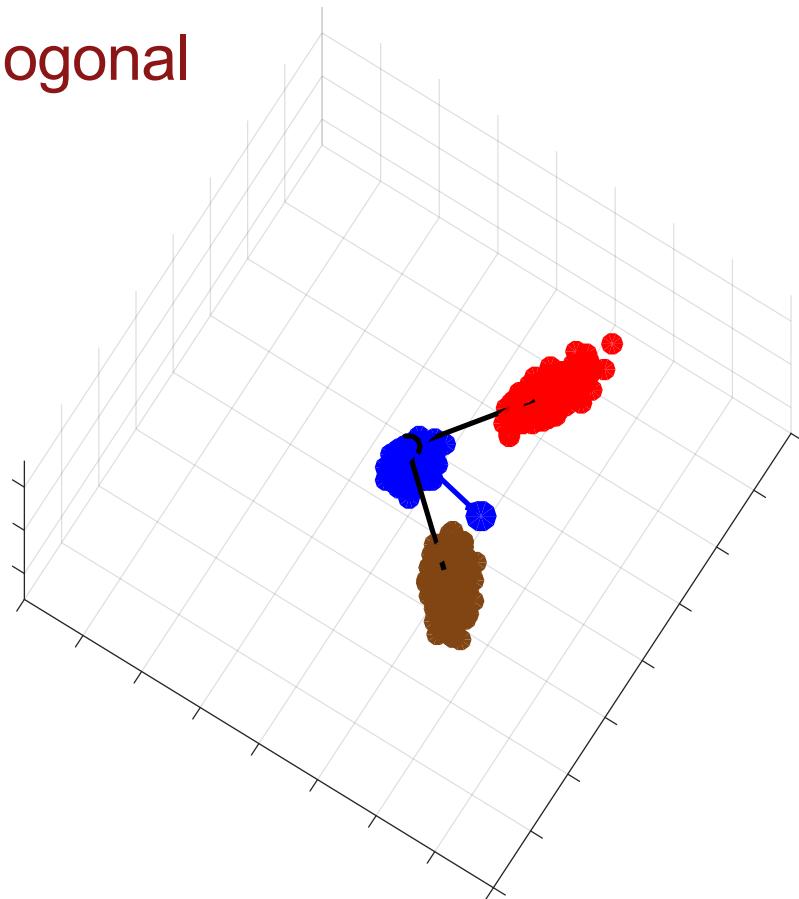
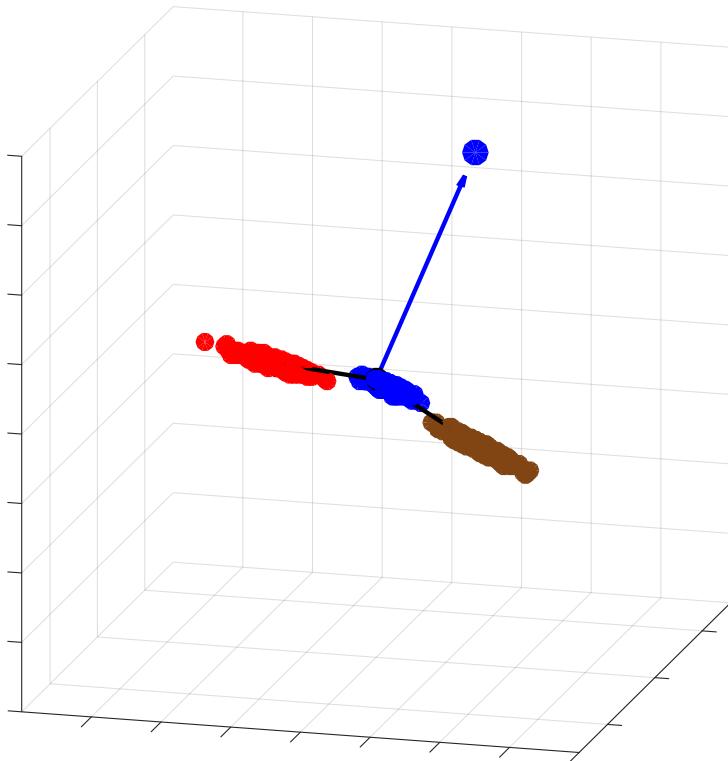
# Distinct clusters are almost orthogonal

- Eigenvector embedding from A:



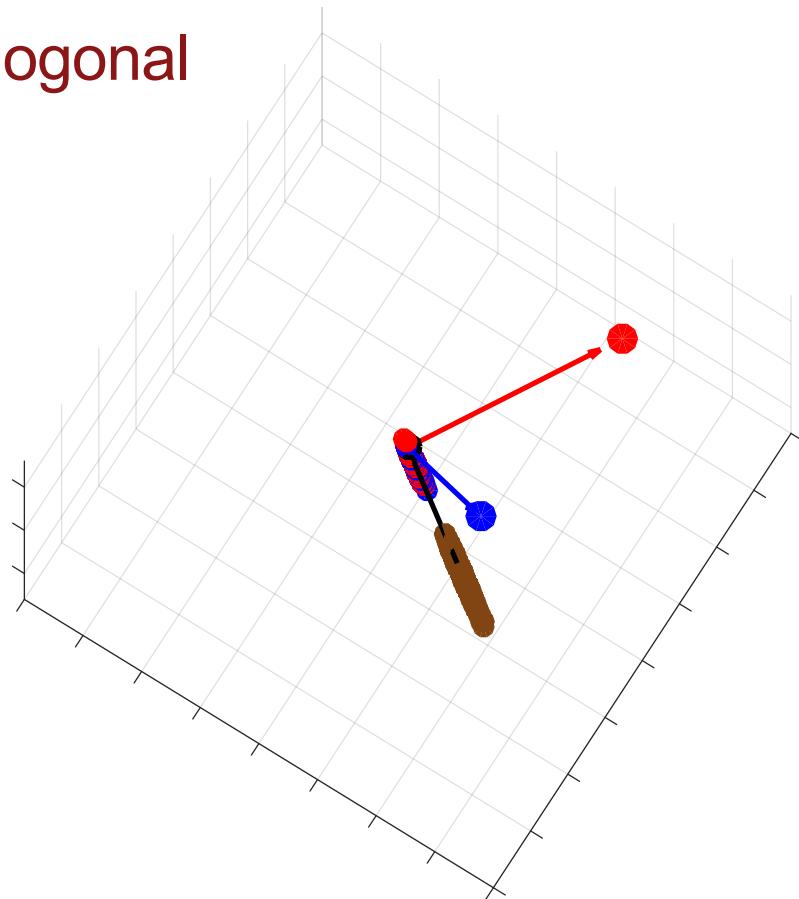
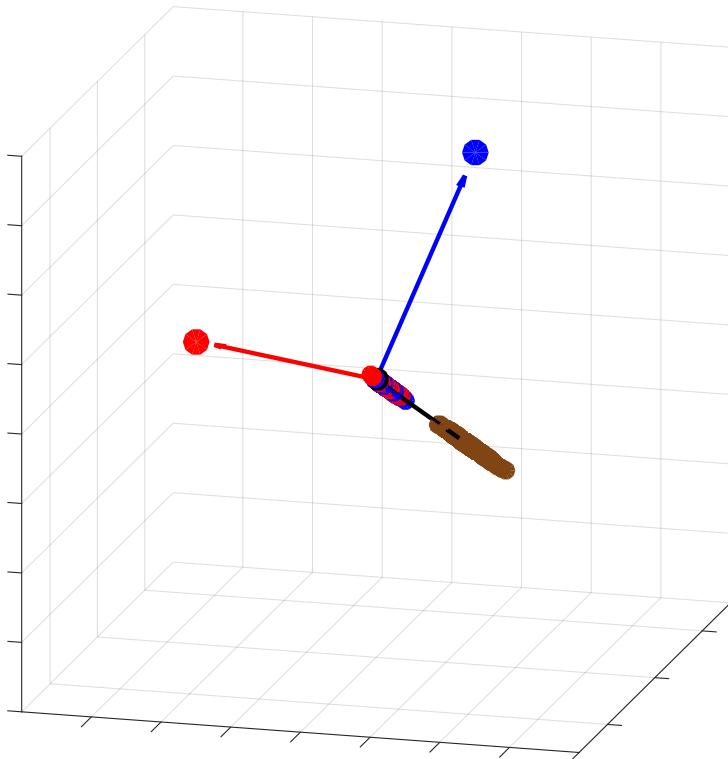
# Distinct clusters are almost orthogonal

- After step 1 of pivoted QR



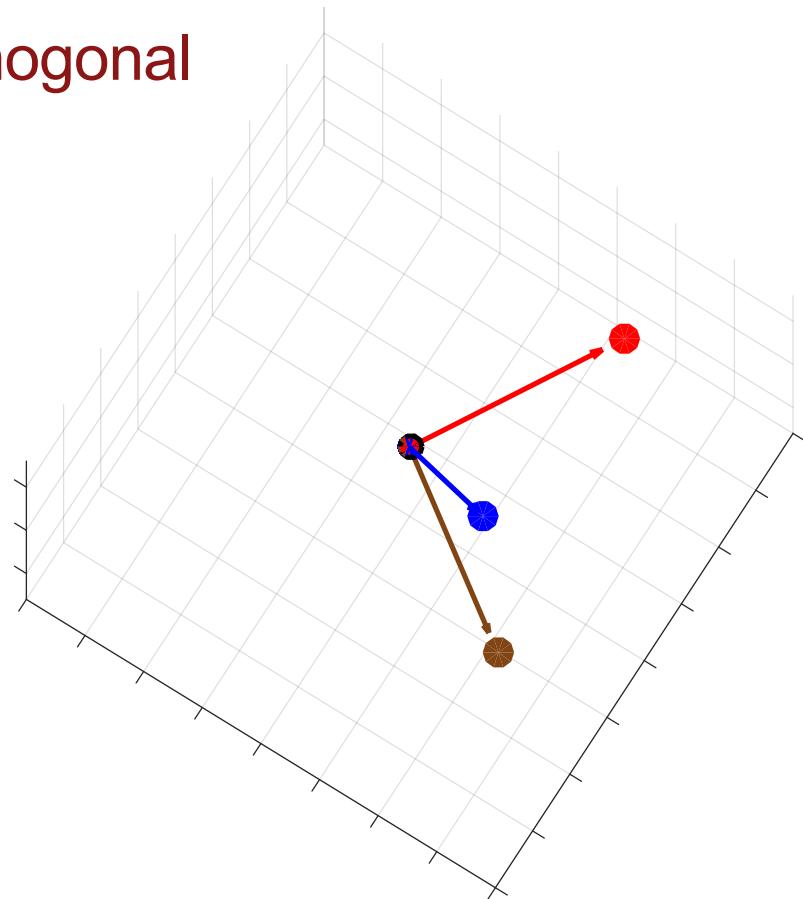
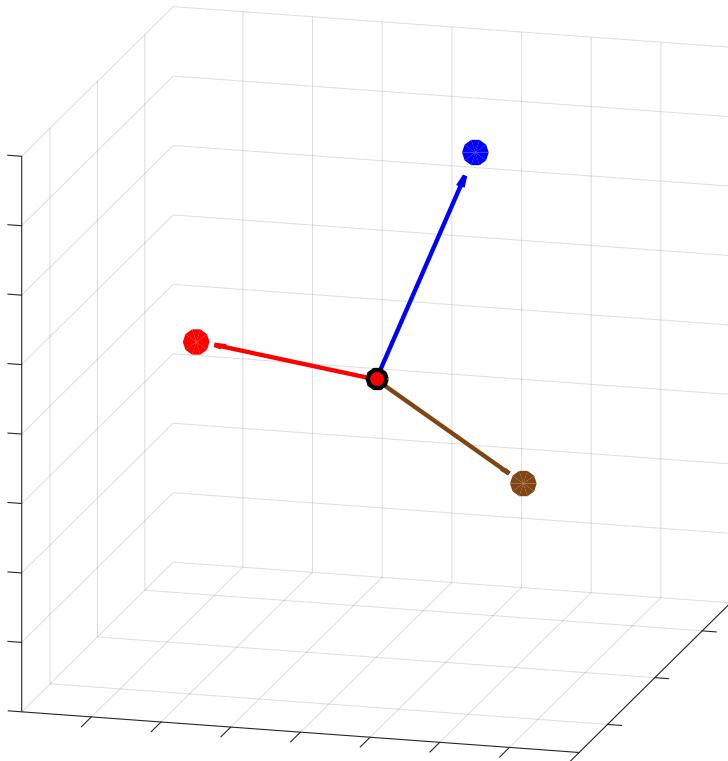
# Distinct clusters are almost orthogonal

- After step 2 of pivoted QR



# Distinct clusters are almost orthogonal

- After step 3 of pivoted QR



## Randomized variant

- Observation: we just need to select a single “good” column from each cluster to get a basis of vectors that are almost indicators.
- Idea: use coupon-collecting argument to subsample the vertices prior to pivoted QR.
- Cost is still  $\mathcal{O}(nk^2)$  overall, but a much smaller constant.

## Our randomized algorithm

1. Sample  $\mathcal{O}(k \log k)$  vertices and denote this set  $\mathcal{J}$ , letting  $W_k = V_k(\mathcal{J}, :)$
2. Compute a column-pivoted QR factorization of  $W_k^T$

$$W_k^T \Pi = QR$$

3. For each  $j = 1, \dots, n$ , assign vertex  $j$  to cluster

$$c_j = \arg \max_i |Q^T V_k^T|_{i,j}$$

---

- Complexity:  $\mathcal{O}(nk) + \mathcal{O}(k^3 \log k) + \mathcal{O}(nk^2)$

# When does it work?

IF eigenvectors of  $A$  are **close enough** to eigenvectors of  $M$  **in span**

$$\|V_k V_k^T - U_k U_k^T\|_2 < \epsilon = o\left(\frac{1}{\sqrt{n}}\right)$$

THEN our (randomized) algorithm recovers the clusters **for sufficiently large  $n$**

---

**Theorem 1.** Let  $W \in \mathbf{R}^{n \times k}$  be a normalized indicator matrix and suppose  $n > 4$ . Let  $V \in \mathbf{R}^{n \times k}$  with orthonormal columns satisfy

$$\|WW^T - VV^T\|_2 \leq \epsilon$$

with  $\epsilon < 2^{-k}/\sqrt{n}$ . If  $V^T \Pi = QR$  is a CPQR factorization of  $V^T$ , then there exists a permutation matrix  $\widehat{\Pi}$  and a diagonal matrix  $D$  with  $D_{ii} = \pm 1$  such that

$$\|W\widehat{\Pi}D - VQ\|_F \leq \epsilon(1 + \sqrt{2})k\sqrt{n} + \mathcal{O}(\epsilon^2).$$

## Back to the stochastic block model

- Identifiability of clusters in SBM exhibits a threshold behavior
  - › See [Abbe et al., 2016], [Agarwal et al., 2015], and [Hajek et al., 2016]

- Parameterize:

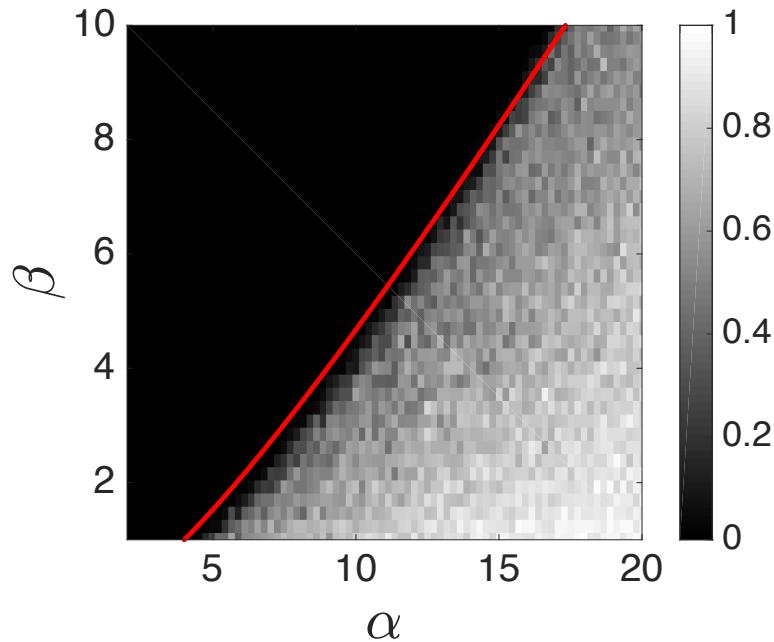
$$\mathbb{P}[A_{ij} = 1] = \begin{cases} \alpha \log m / m & i \text{ and } j \text{ in same cluster,} \\ \beta \log m / m & \text{else.} \end{cases}$$

- Result: asymptotically in number of nodes, recovery is possible with high probability if and only if

$$\sqrt{\alpha} - \sqrt{\beta} > 1$$

# So, what's the problem?

- Spectral clustering with k-means++ on eigenvector embedding
  - › Bad recovery results (right)
  - › Color is empirical recovery probability
  - › Need a good initialization to avoid local minima!
  
- A semidefinite program over  $n$ -by- $n$  semidefinite matrices attains asymptotically optimal recovery
  - › [Abbe et al., 2016], [Agarwal et al., 2015], and [Hajek et al., 2016]
  - › Too expensive for large  $n$

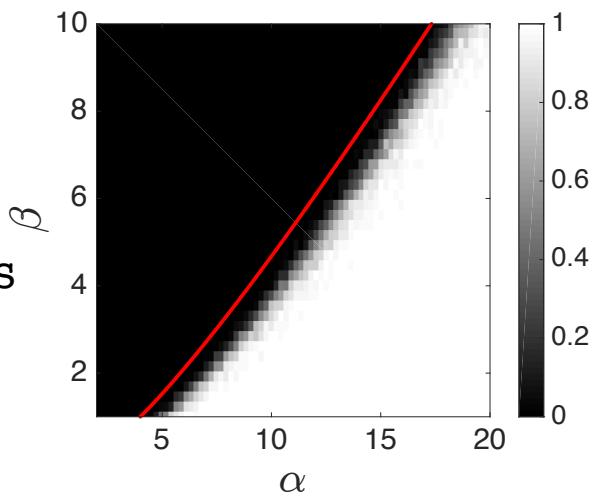


# Results: equisized clusters

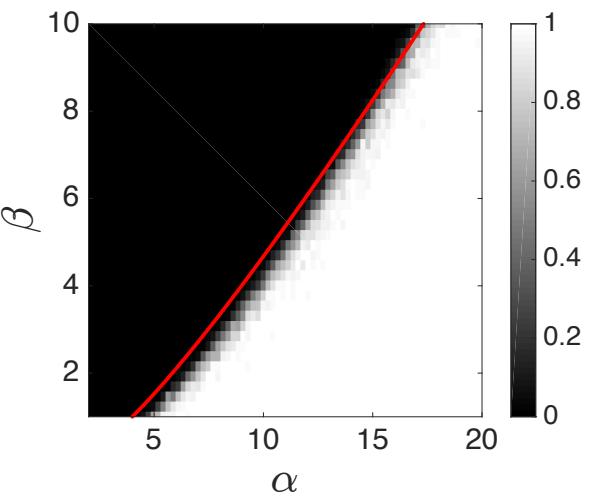
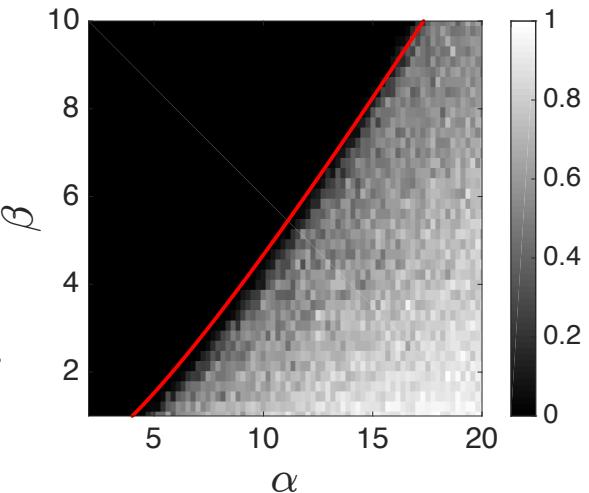
- $n = 1350$  nodes
- $k = 9$  clusters

Left: QR

Right: QR + k-means



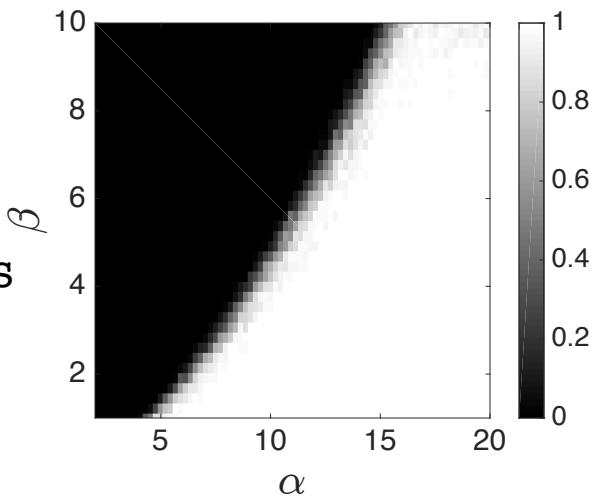
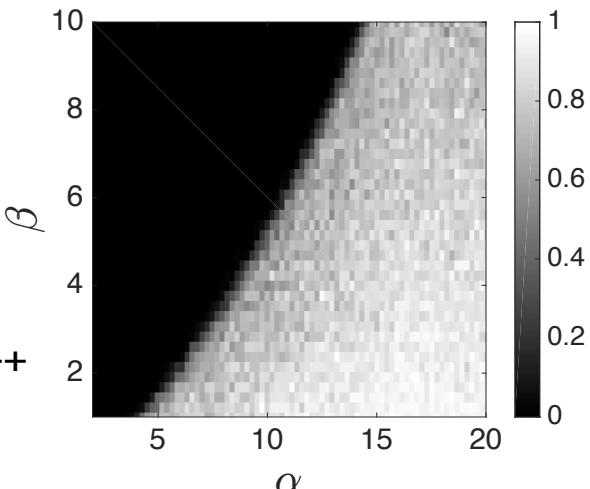
Right: k-means++



# Results: unequisized clusters

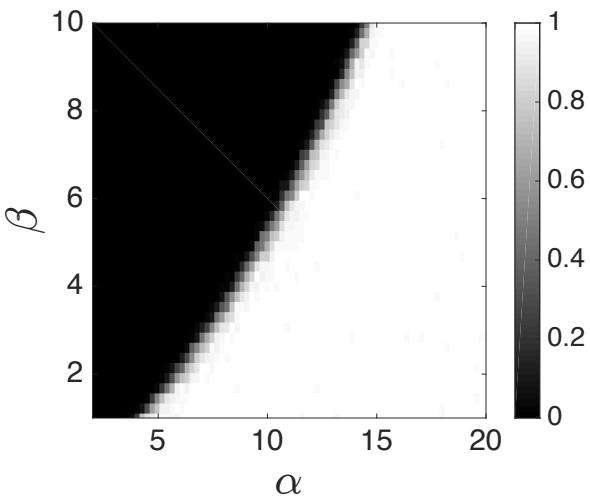
- $m \in [70,130]$  nodes/cluster
- $k = 7$  clusters

Recall: k-means++



Left: QR

Right: QR + k-means



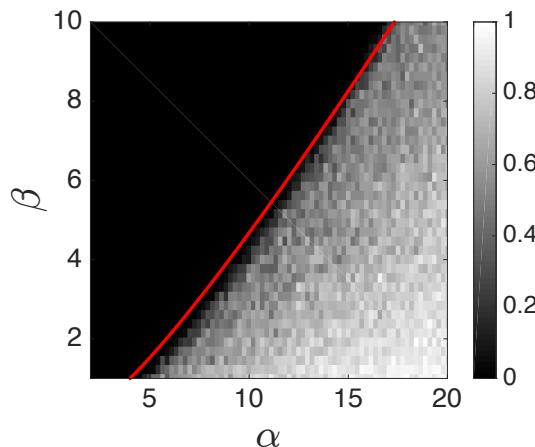
Stanford University

## Related work for cluster recovery of sparse SBM

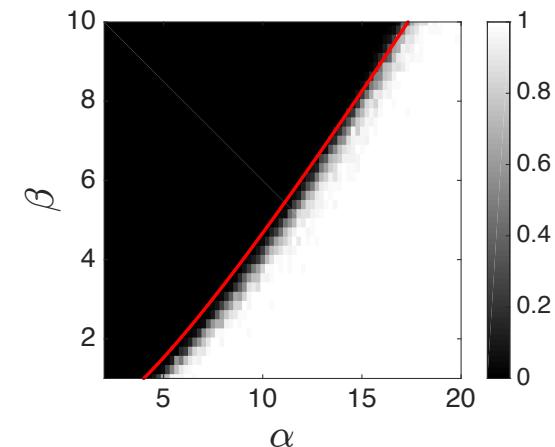
- [Zha et al., 2001]: a similar pivoted QR based algorithm for general data
  - › Different cluster assignment scheme
  - › Not in the context of graphs
  - › Deterministic only
- [Bopanna, 1987] and [McSherry, 2001]: more complicated spectral-based algorithms, with **provable recovery guarantees** but not at the limit
- [Mossel et al., 2016]: spectral algorithm plus post-processing for  $k=2$ , with recovery guarantees at correct limit

# Ongoing and future work

- Can we prove asymptotically optimal recovery with this algorithm (or a variant)?
- Application to optimization problems in computer vision
- Take-away:

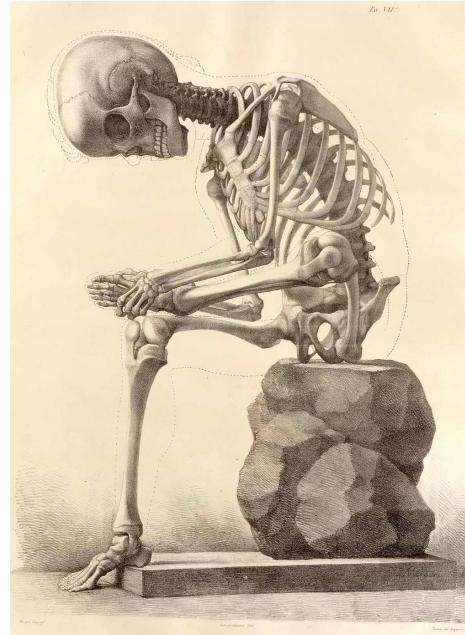


$$[Q, R, P] = qr(Vk', 0);$$
$$[\tilde{c}, c] = \max(\text{abs}(R^*P'));$$



# A recursive skeletonization factorization based on strong admissibility

WITH HO, DAMLE, YING

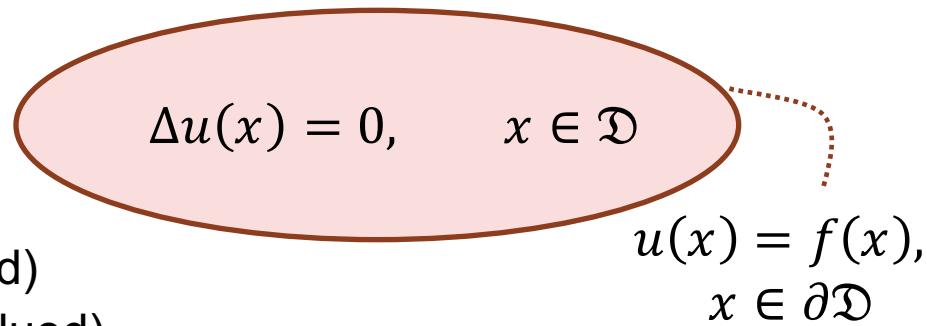


# Introduction

- Many elliptic differential equations describing physical phenomena can be recast as integral equations of the form

$$a(x)u(x) + \int_{\Omega} b(x)K(x-y)c(y)u(y) dy = f(x), \quad x \in \Omega$$

- Examples
  - › Laplace equation
  - › Helmholtz equation
  - › Lippmann-Schwinger
  - › Stokes equation (vector-valued)
  - › Elasticity equations (vector-valued)



# Introduction

- Let's simplify:

$$\int_{\Omega} K(x - y)u(y) dy = f(x), \quad x \in \Omega$$

- $K(x - y) = K(r) = \dots$ 
  - ›  $1/|r|$
  - ›  $\log(|r|)$
  - › A derivative of one of these

$$\Delta u(x) = 0, \quad x \in \mathfrak{D}$$
$$u(x) = f(x), \quad x \in \partial \mathfrak{D}$$

## Discretization

- Discretization using standard methods leads to a many-body sum (assume uniform weights)

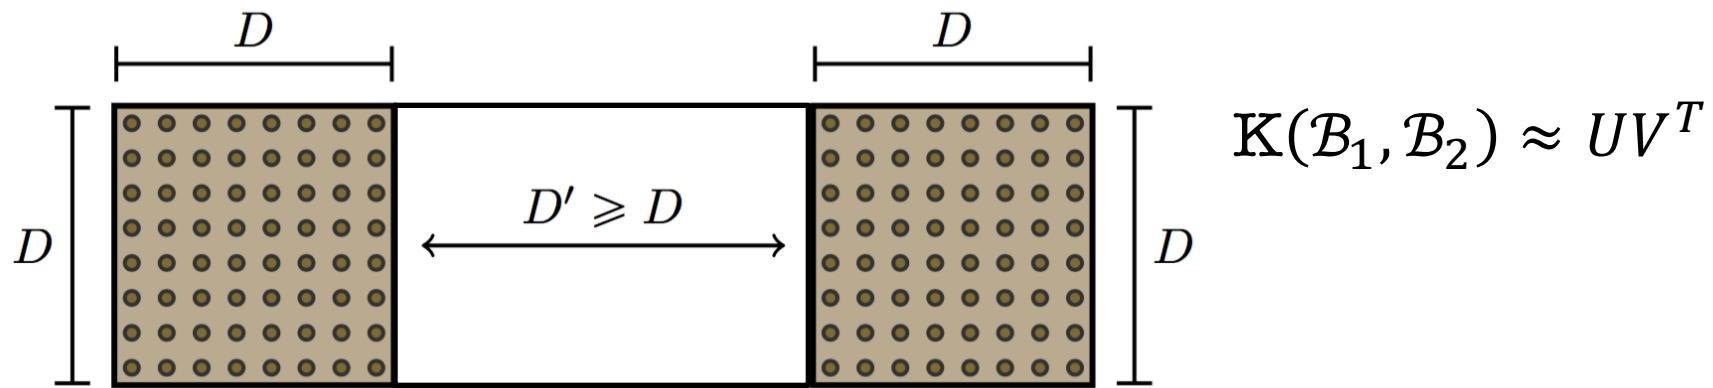
$$\sum_j K(x_i - x_j) u_j = f_i, \quad i = 1, \dots, N$$

- Naïvely, evaluating the output  $f$  given the input  $u$  costs  $\mathcal{O}(N^2)$
- Tree codes [Barnes & Hut, 1986] or fast multipole method (FMM) [Greengard & Rokhlin, 1987 & 1997] use clever approximations to reduce this to  $\mathcal{O}(N \log N)$  or  $\mathcal{O}(N)$  with a constant depending on accuracy.

# Tree codes and FMM

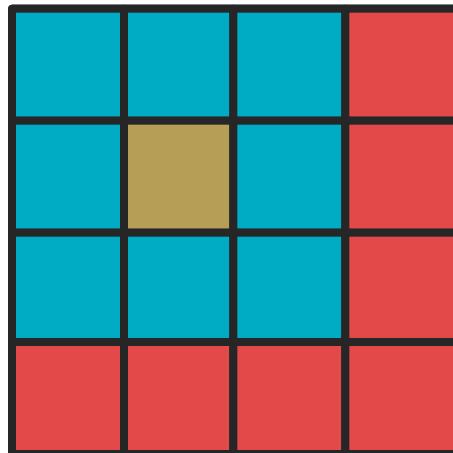
- Basic idea: well-separated interactions are compressible

$$\sum_j K(x_i - x_j) u_j = f_i, \quad i = 1, \dots, N$$

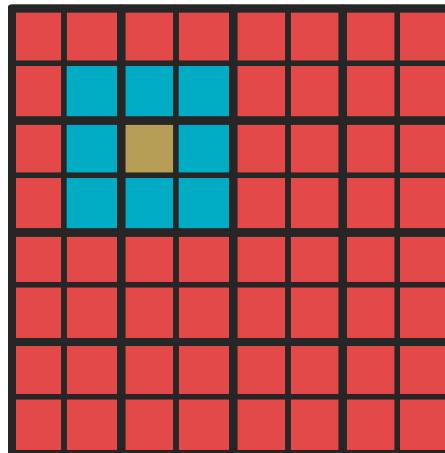


# Tree codes and FMM

- Exploit far-field compression idea in a **hierarchical fashion** to apply operator quickly
- This is as much as we need to understand tree codes for this talk



One level of tree



Next lower level of tree

# Solving systems

- What about the opposite problem, solving  $Ku = f$  for  $u$ ?

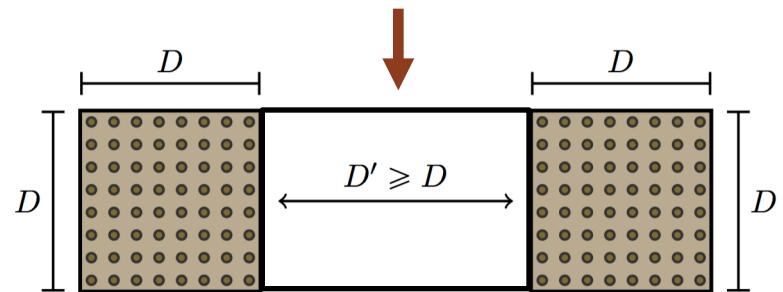
$$\sum_j K(x_i - x_j)u_j = f_i, \quad i = 1, \dots, N$$

- › Iterative methods (conjugate gradient, etc.) are  $\mathcal{O}(N)$  per iteration with FMM, but often require many iterations.
- Large body of work on approximate solvers or preconditioners using hierarchical low-rank structure of this problem, giving better complexity.

# Hierarchical representations from skeletonization

- “Recursive skeletonization” and related literature has led to many nice hierarchical factorizations based on “weak admissibility”
  
- HSS / HBS matrices
  - › [Martinsson & Rokhlin, 2005]
  - › [Chandrasekaran et al., 2006 & 2007]
  - › [Ho & Greengard, 2012]
  - › [Xia et al., 2012]
  - › [Gillman et al., 2012]
  
- HODLR matrices
  - › [Martinsson, 2008]
  - › [Ambikasaran & Darve, 2013]

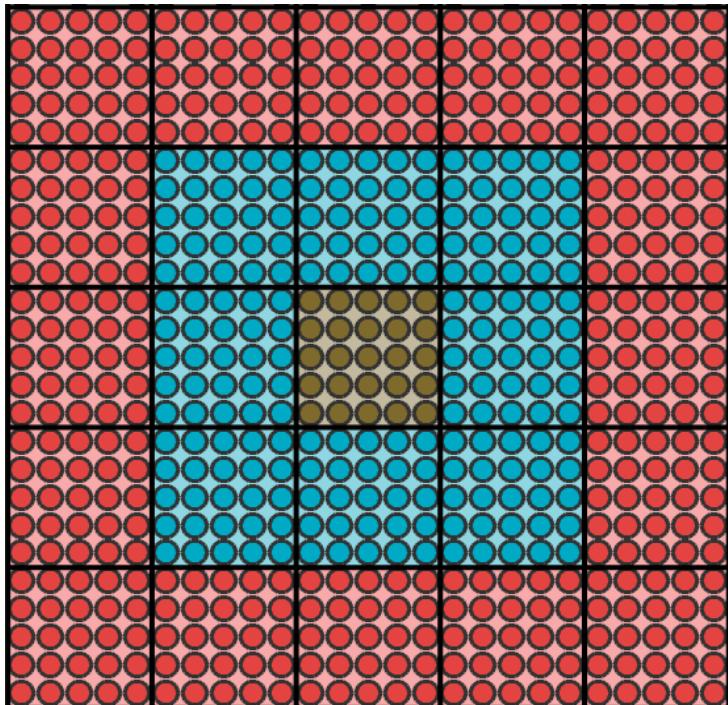
Separation not enforced  
by HSS/HBS/HODLR,  
giving less compression



# Our approach

- Adapt recursive skeletonization [Martinsson & Rokhlin, 2005] to compress only well-separated interactions using the multiplicative formulation of skeletonization in [Ho & Ying, 2016].
- Result: approximate factorization of  $K^{-1}$  as the product of
  - › permutation matrices
  - › block-unit-triangular matrices (with small blocks)
  - › a block-diagonal matrix (with small blocks)
- Simple and fast way to apply  $K$ ,  $K^{-1}$ ,  $K^{1/2}$ , or  $K^{-1/2}$ , or compute log-det

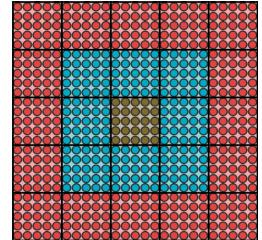
# Quadtree decomposition and admissible neighbors



- **Brown:** box b
- **Blue:** near-field neighbors of b
- **Red:** far-field neighbors of b  
(lots of these)

$$\begin{bmatrix} A_{bb} & A_{bn} & A_{bf} \\ \hline A_{nb} & A_{nn} & A_{nf} \\ \hline A_{fb} & A_{fn} & A_{ff} \end{bmatrix}$$

# Interpolative decomposition



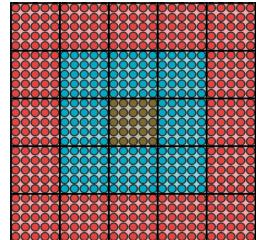
- By assumption, the interactions between **box  $b$**  and **far-field neighbors  $f$**  are low-rank.
- Compress these with an *interpolative decomposition* [Cheng et al., 2005], where box  **$b$**  is partitioned into “**skeleton set**”  **$s$**  and “**redundant set**”  **$r$**

$$\begin{bmatrix} A_{bb} & A_{bn} & A_{bf} \\ \hline A_{nb} & A_{nn} & A_{nf} \\ \hline A_{fb} & A_{fn} & A_{ff} \end{bmatrix}$$

$$b = s \cup r$$

$$A_{fr} \approx A_{fs} T$$

## Block sparse elimination

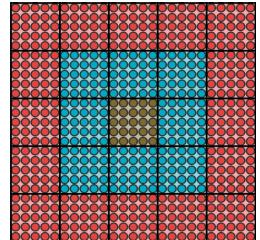


$$\left[ \begin{array}{c|c|c} A_{bb} & A_{bn} & A_{bf} \\ \hline A_{nb} & A_{nn} & A_{nf} \\ \hline A_{fb} & A_{fn} & A_{ff} \end{array} \right] = \left[ \begin{array}{c|c|c|c} A_{rr} & A_{rs} & A_{rn} & A_{rf} \\ \hline A_{sr} & A_{ss} & A_{sn} & A_{sf} \\ \hline A_{nr} & A_{ns} & A_{nn} & A_{nf} \\ \hline A_{fr} & A_{fs} & A_{fn} & A_{ff} \end{array} \right]$$

$$b = s \cup r$$

$$A_{fr} \approx A_{fs} T$$

# Block sparse elimination

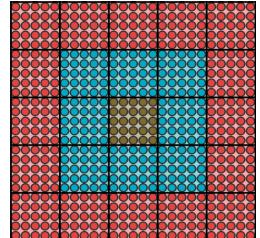


$$\left[ \begin{array}{c|c|c} A_{bb} & A_{bn} & A_{bf} \\ \hline A_{nb} & A_{nn} & A_{nf} \\ \hline A_{fb} & A_{fn} & A_{ff} \end{array} \right] \approx \left[ \begin{array}{c|c|c|c} A_{rr} & A_{rs} & A_{rn} & T^T A_{sf} \\ \hline A_{sr} & A_{ss} & A_{sn} & A_{sf} \\ \hline A_{nr} & A_{ns} & A_{nn} & A_{nf} \\ \hline A_{fs} T & A_{fs} & A_{fn} & A_{ff} \end{array} \right]$$

$$b = s \cup r$$

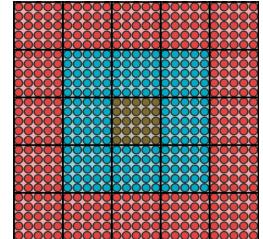
$$A_{fr} \approx A_{fs} T$$

## Block sparse elimination



$$\begin{array}{c|c|c|c} \begin{matrix} A_{rr} & A_{rs} & A_{rn} & T^T A_{sf} \\ \hline A_{sr} & A_{ss} & A_{sn} & A_{sf} \\ \hline A_{nr} & A_{ns} & A_{nn} & A_{nf} \\ \hline A_{fs} T & A_{fs} & A_{fn} & A_{ff} \end{matrix} & \longrightarrow & \begin{matrix} X_{rr} & X_{rs} & X_{rn} & \\ \hline X_{sr} & A_{ss} & A_{sn} & A_{sf} \\ \hline X_{nr} & A_{ns} & A_{nn} & A_{nf} \\ \hline & A_{fs} & A_{fn} & A_{ff} \end{matrix} \end{array}$$

## Block sparse elimination

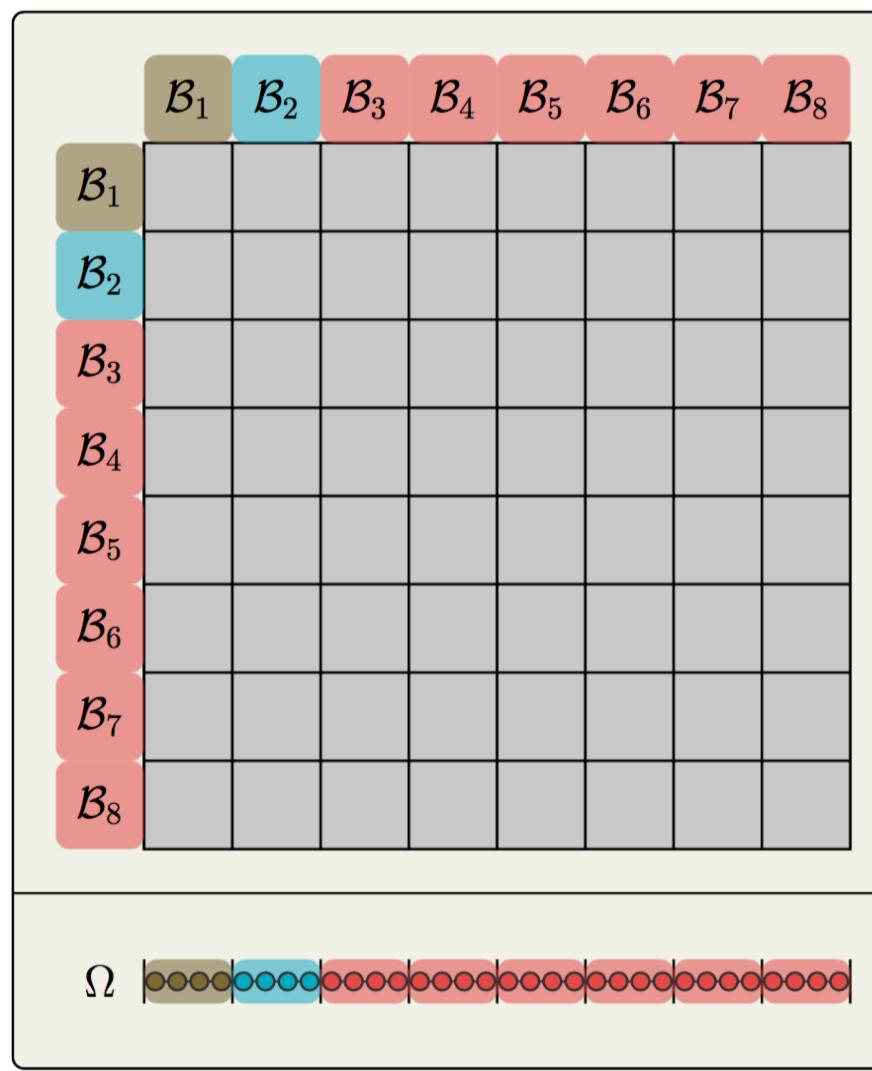


$$\begin{array}{cc|cc|c} X_{rr} & X_{rs} & X_{rn} & & \\ X_{sr} & A_{ss} & A_{sn} & A_{sf} & \\ \hline X_{nr} & A_{ns} & A_{nn} & A_{nf} & \\ \hline & A_{fs} & A_{fn} & A_{ff} & \end{array} \longrightarrow \begin{array}{c|c|c|c} X_{rr} & & & \\ & X_{ss} & X_{sn} & A_{sf} \\ \hline & X_{ns} & X_{nn} & A_{nf} \\ \hline & A_{fs} & A_{fn} & A_{ff} \end{array}$$

After block Gaussian elimination, we have decoupled the redundant points and updated interactions between the skeleton points and near-field neighbors.

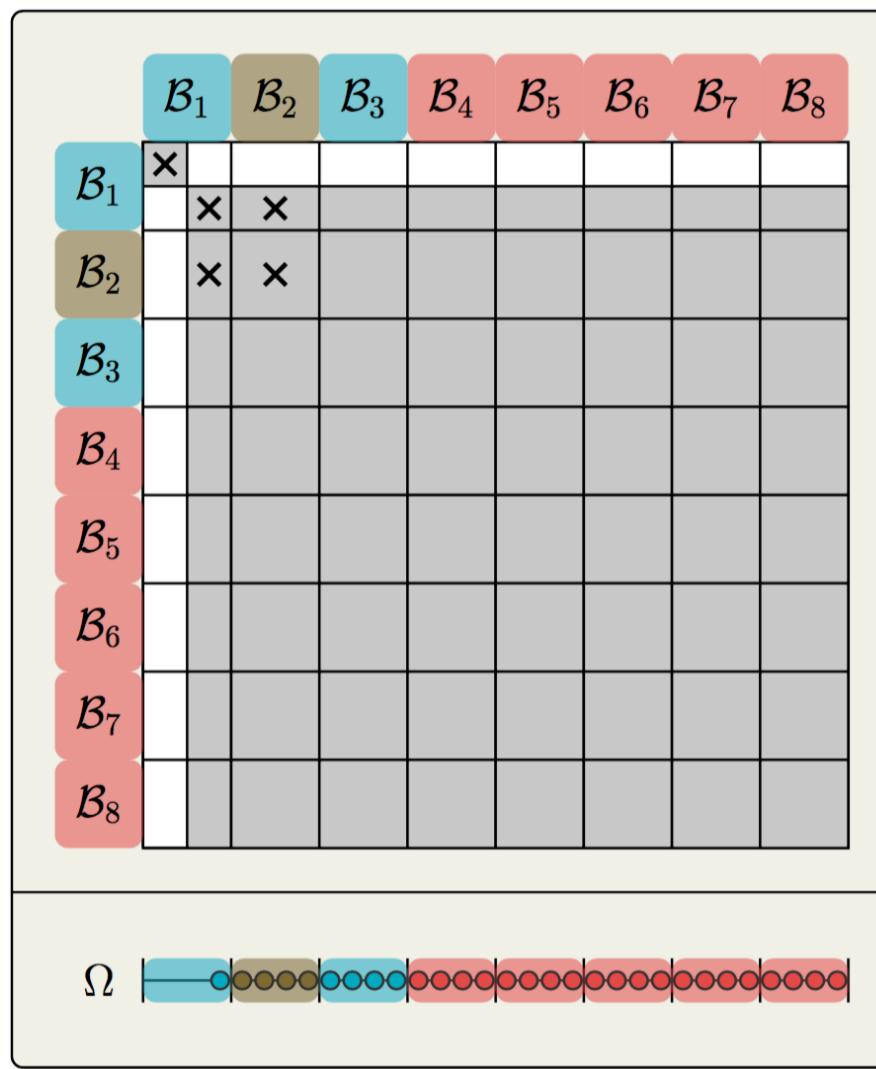
Note: every operation was cheap and is easy to undo

# 1D Level 1 Before $\mathcal{B}_1$



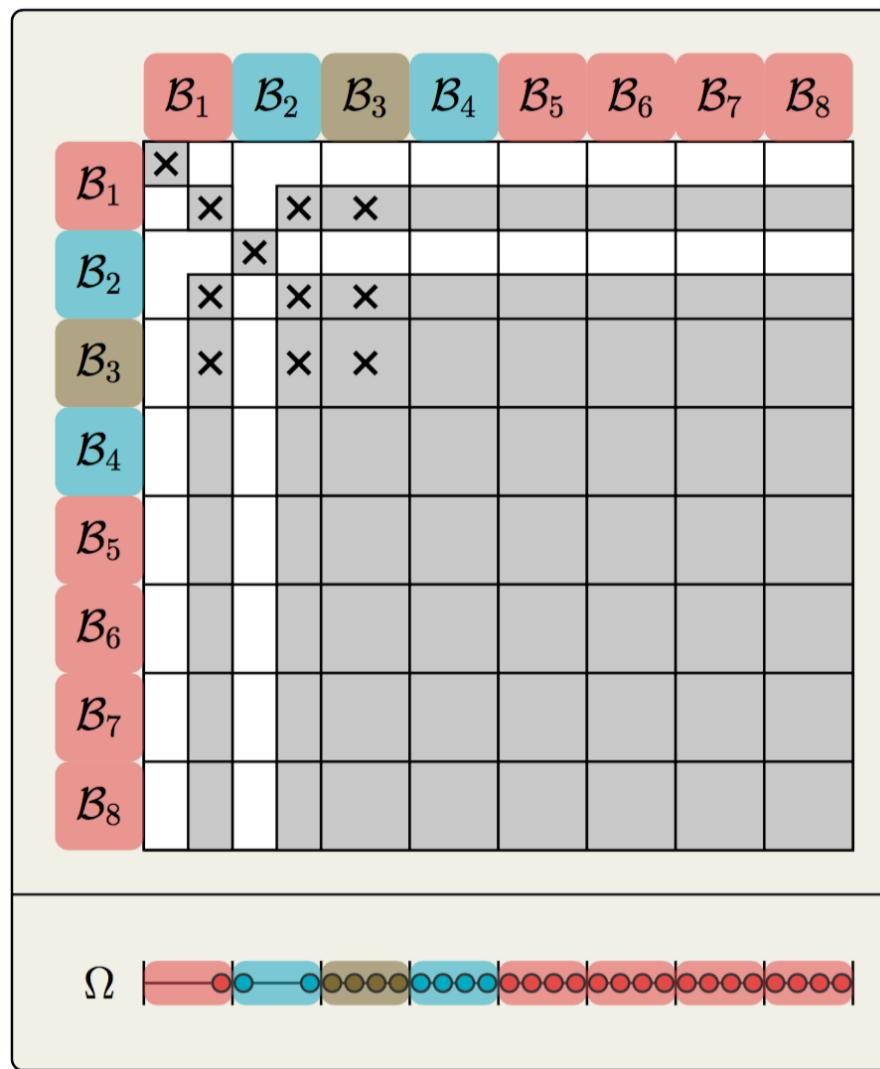
1D  
Level 1  
Before  $\mathcal{B}_2$

Points in  
the domain



← The matrix

1D  
Level 1  
Before  $\mathcal{B}_3$

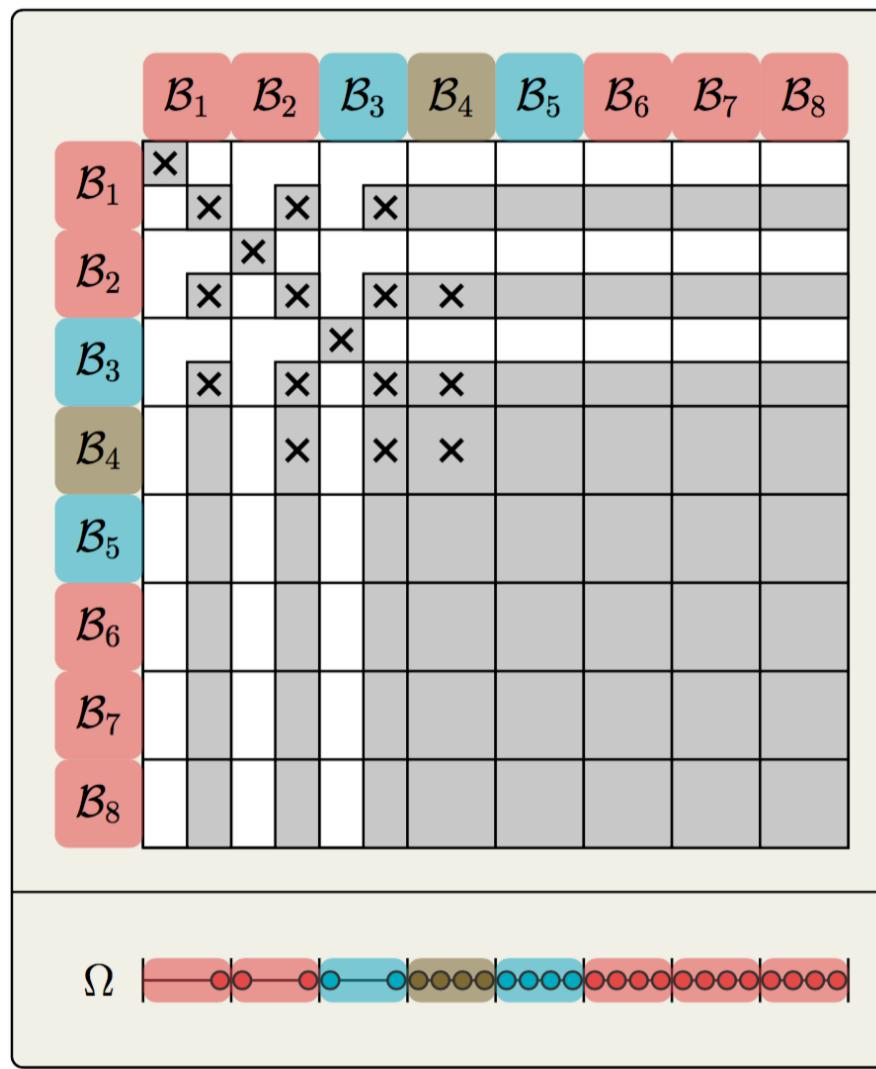


← The matrix

Points in  
the domain

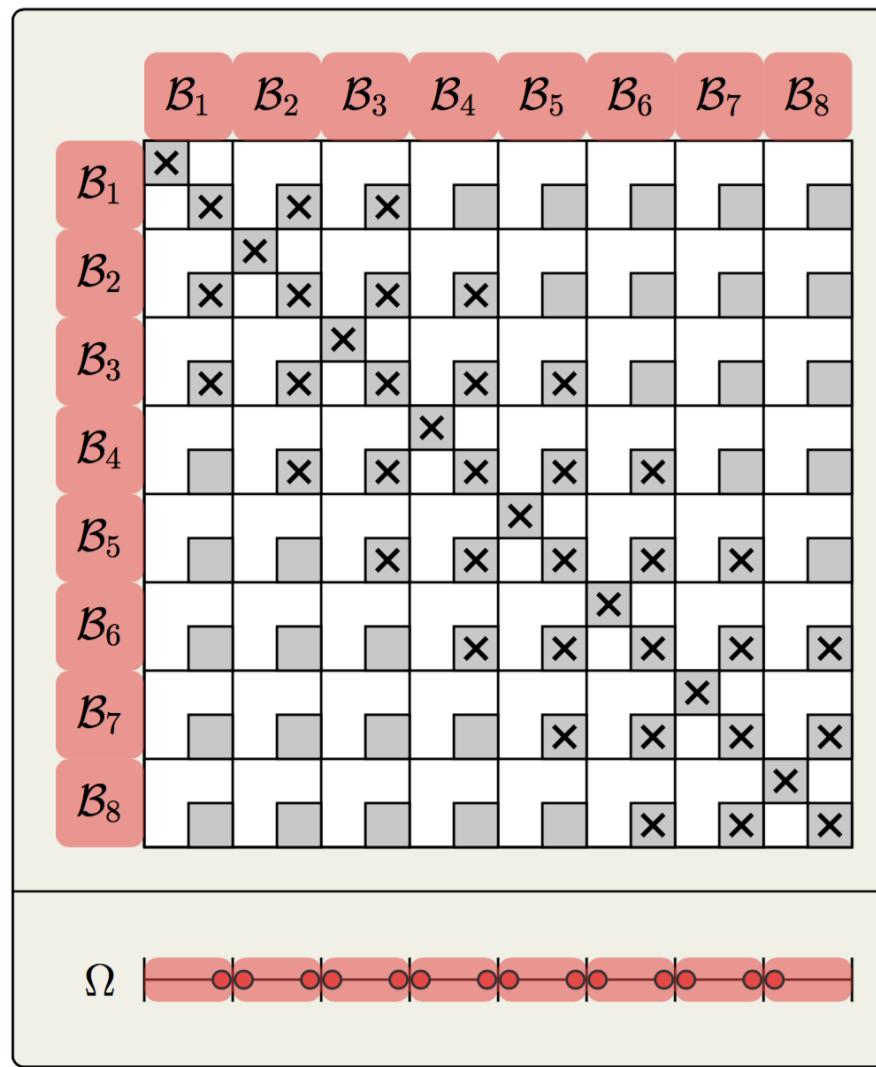
1D  
Level 1  
Before  $\mathcal{B}_4$

Points in  
the domain

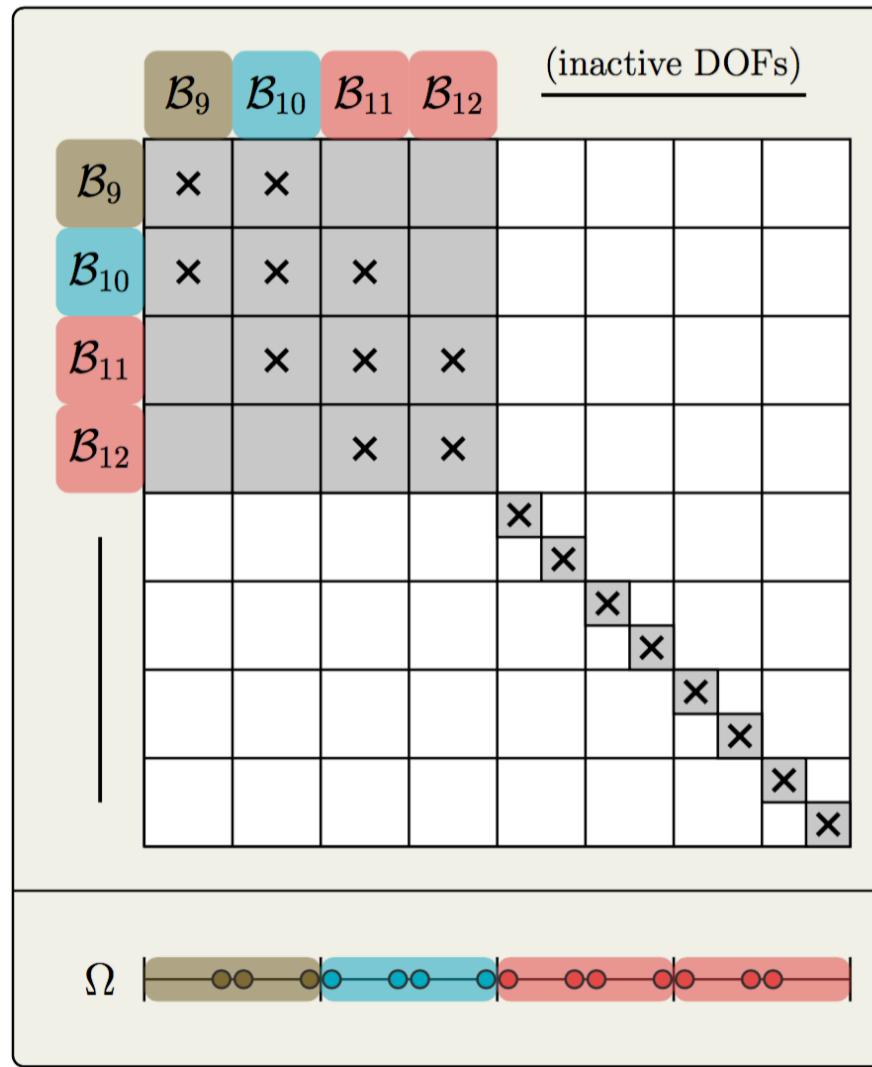


# 1D Level 1 After $\mathcal{B}_8$

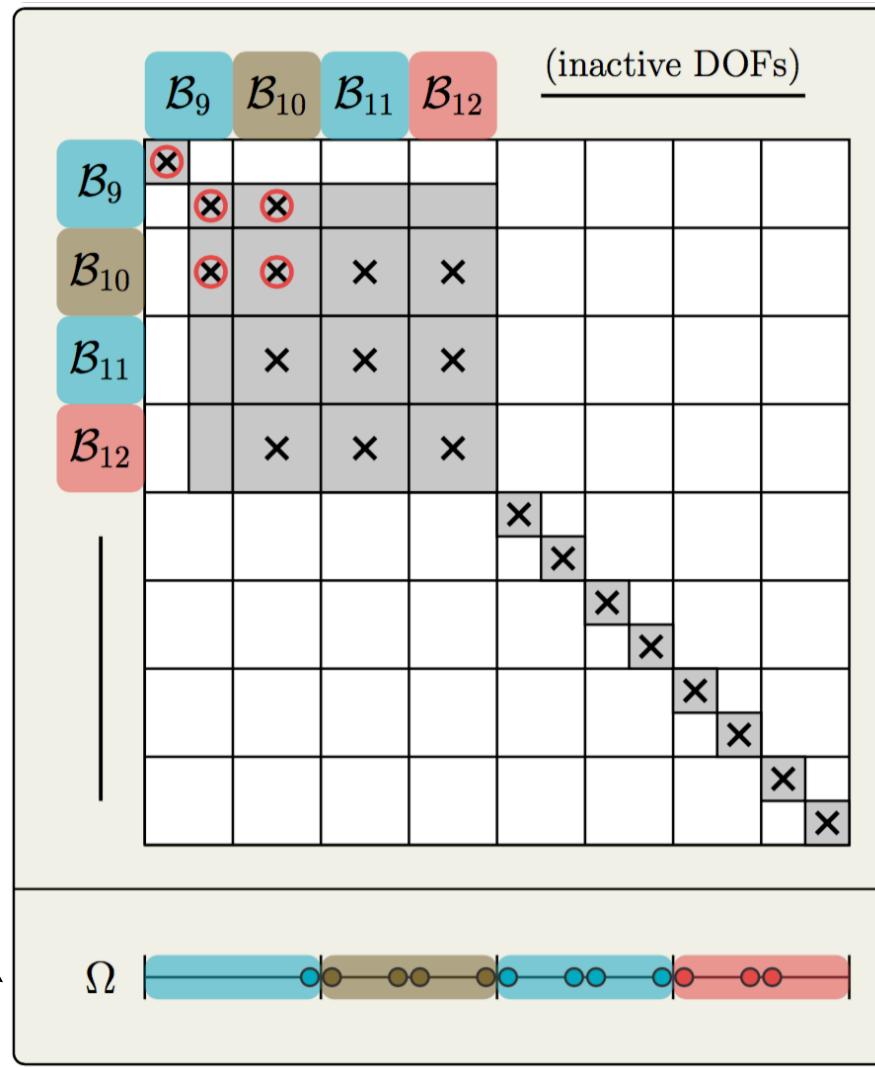
Points in  
the domain



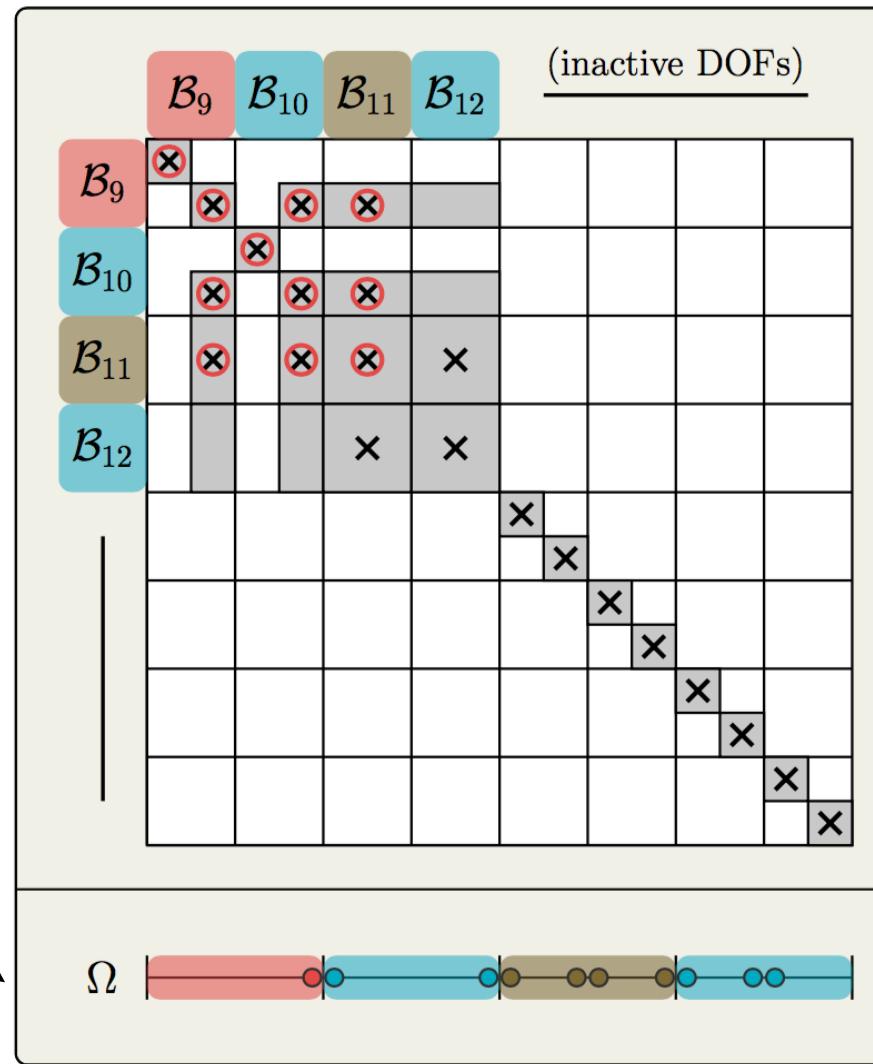
1D  
Level 2  
Before  $\mathcal{B}_9$



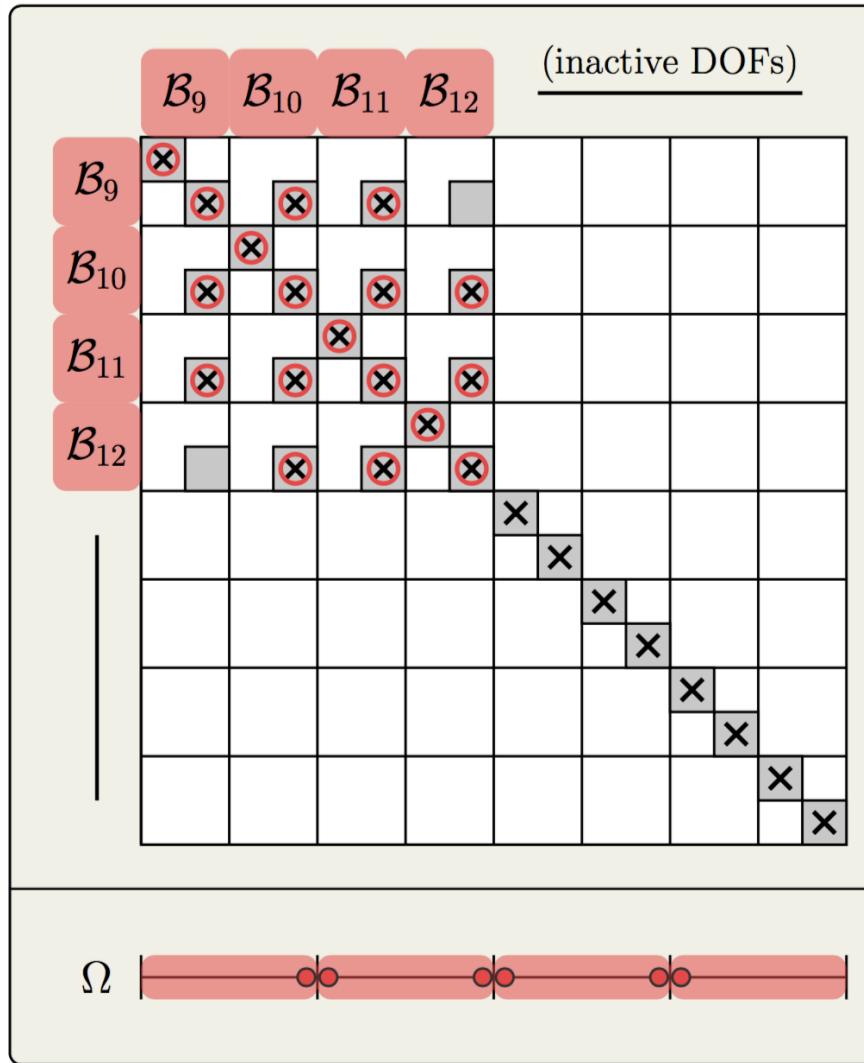
# 1D Level 2 Before $\mathcal{B}_{10}$



1D  
Level 2  
Before  $\mathcal{B}_{11}$

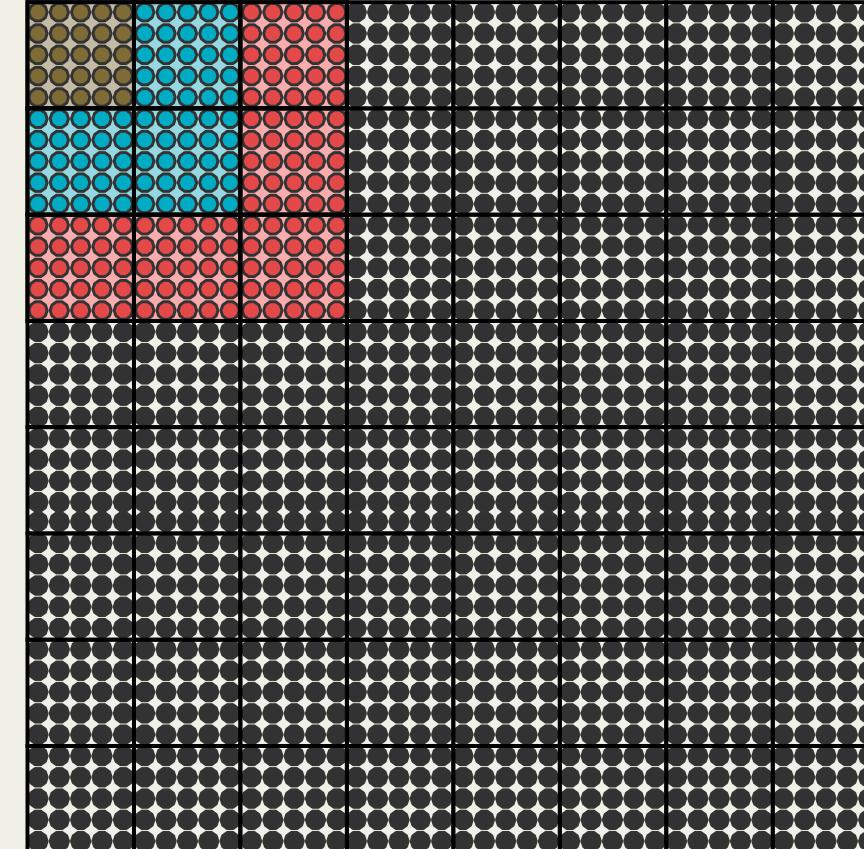


# 1D Level 2 After $\mathcal{B}_{12}$

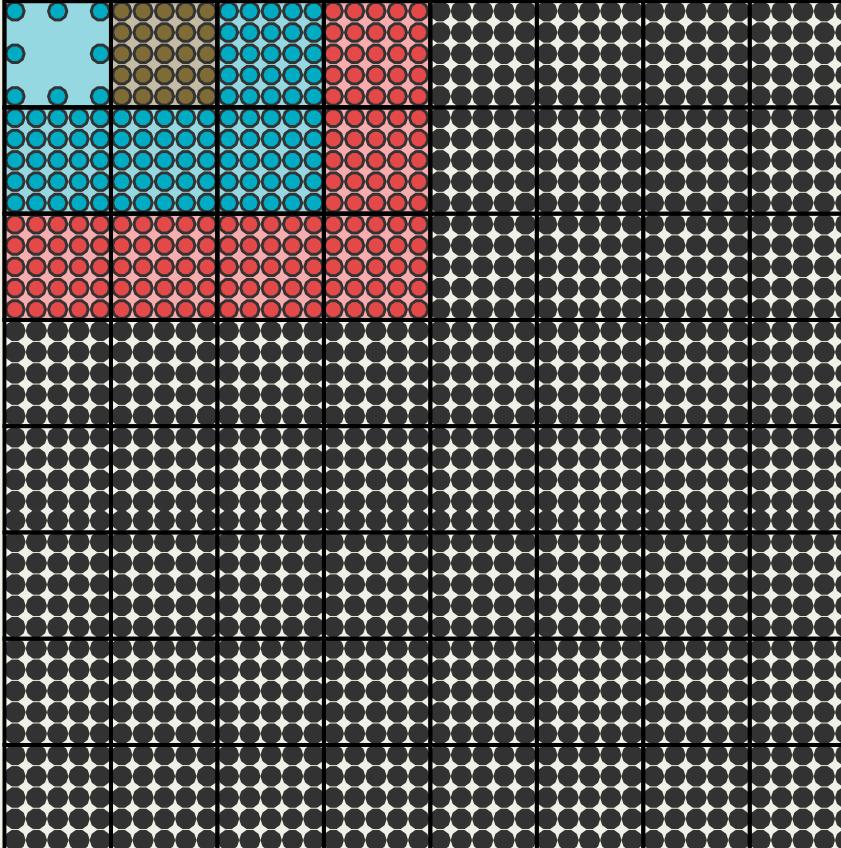


Next: finish  $\mathcal{B}_{12}$  and  
invert (or factor)  
diagonal blocks

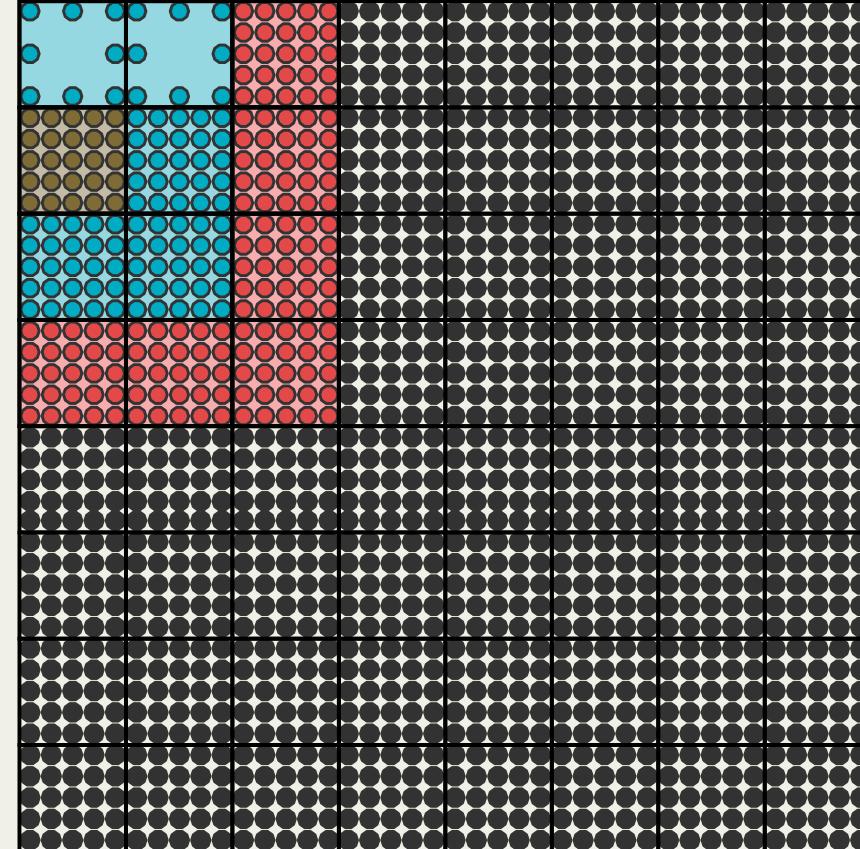
2D  
Level 1  
Before  $\mathcal{B}_1$



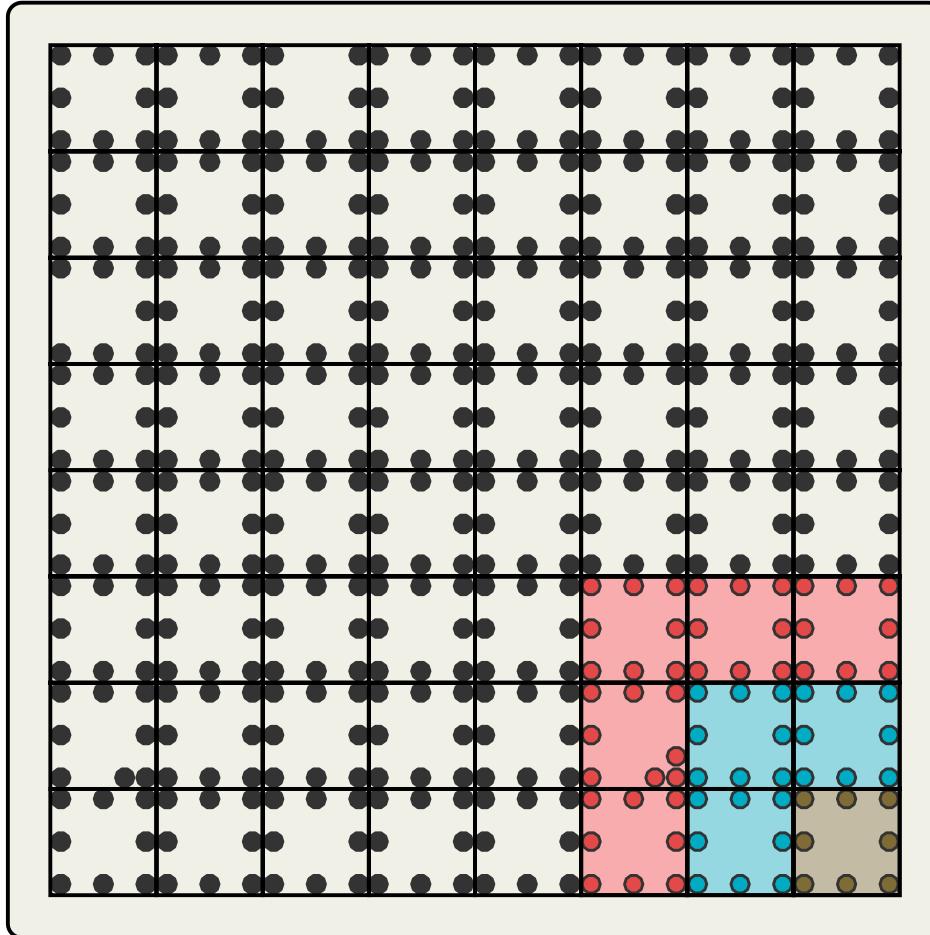
2D  
Level 1  
Before  $\mathcal{B}_2$



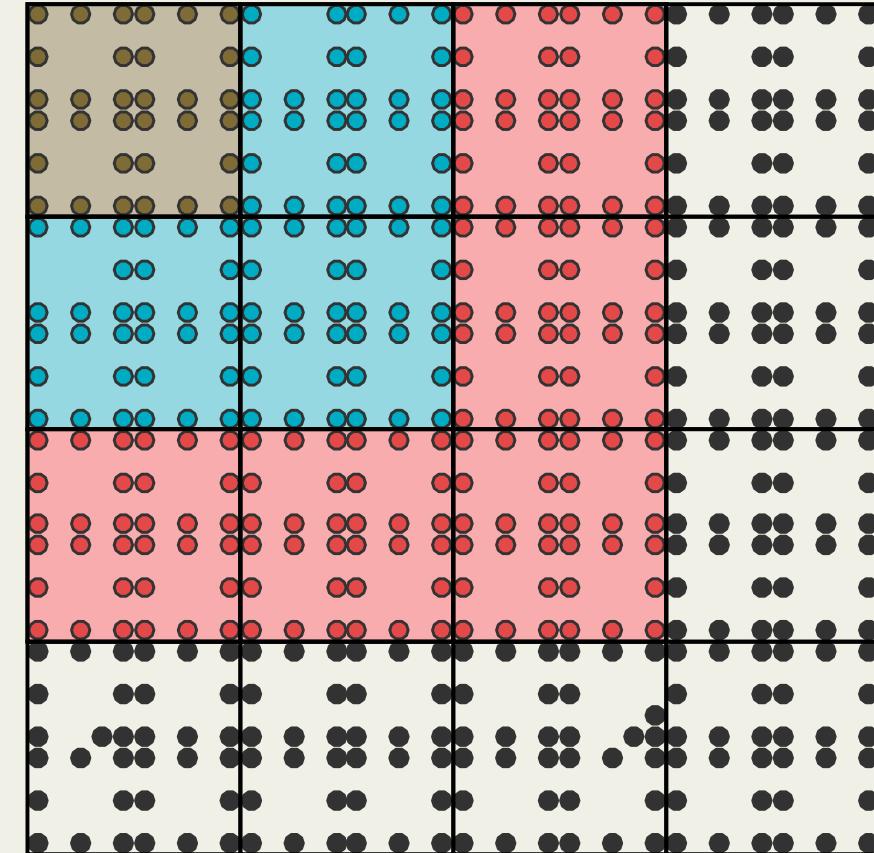
2D  
Level 1  
Before  $\mathcal{B}_3$



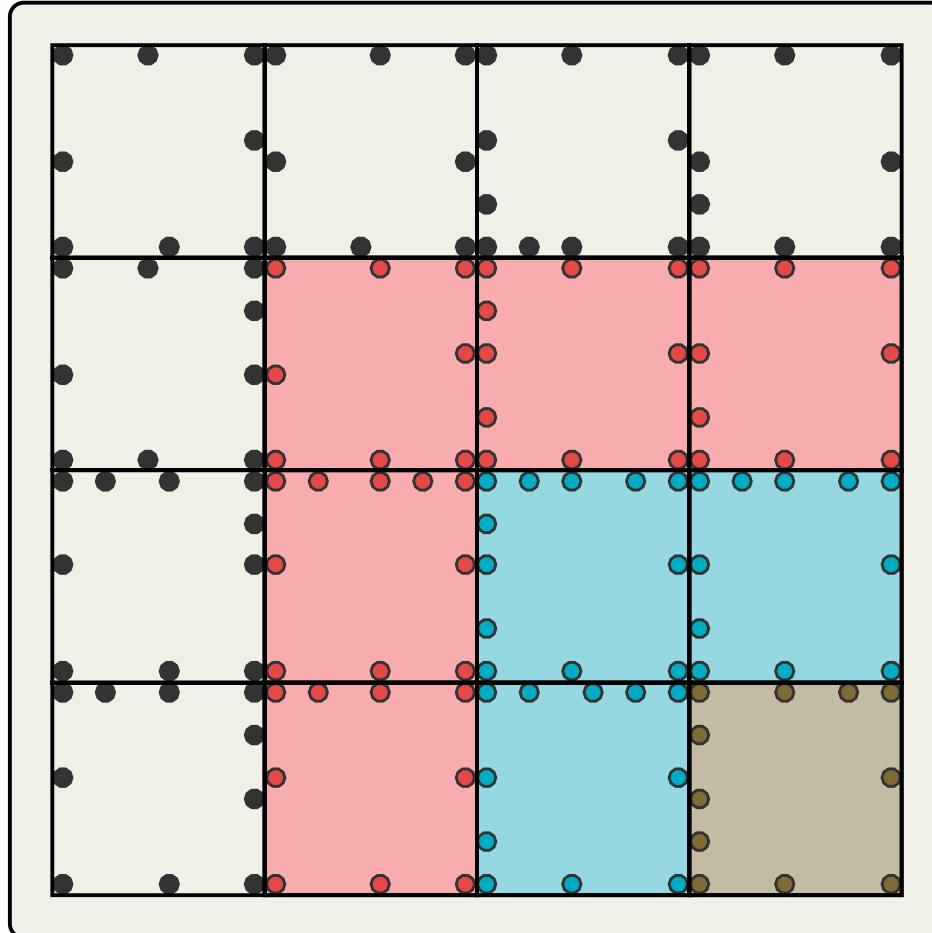
2D  
Level 1  
After  $\mathcal{B}_{64}$



2D  
Level 2  
Before  $B_{65}$



2D  
Level 2  
After  $\mathcal{B}_{80}$

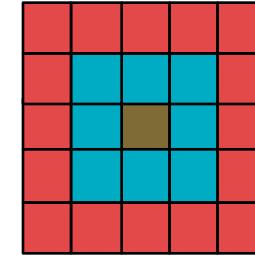


# Complexity sketch

- For each box:
  1. Compute interpolative decomposition
  2. Perform block elimination
- If proxy surface used and ranks do not grow in N:

$\mathcal{O}(N)$  (linear complexity)

- Note: hard part is constant factors  
(store/retrieve Schur updates efficiently)



$$b = s \cup r$$
$$A_{fr} \approx A_{fs} T$$

$$\begin{bmatrix} X_{rr} & & & \\ \hline & X_{ss} & X_{sn} & A_{sf} \\ \hline & X_{ns} & X_{nn} & A_{nf} \\ \hline & A_{fs} & A_{fn} & A_{ff} \end{bmatrix}$$

# Related approaches

- $\mathcal{H}$ -matrices
  - › [Hackbusch & collaborators, 1999-2002]
  - › [Bebendorf, 2008]
- [Coulier et al., 2015] and [Ambikasaran & Darve, 2014] also look at applying an approximate inverse FMM using strong admissibility
  - › Compared to our approach:
    - Multiplicative factorization versus additive decomposition
    - Skeletonization versus more general framework
    - We have extension to hybrid admissibility

## Some scaling results (3D unit cube)

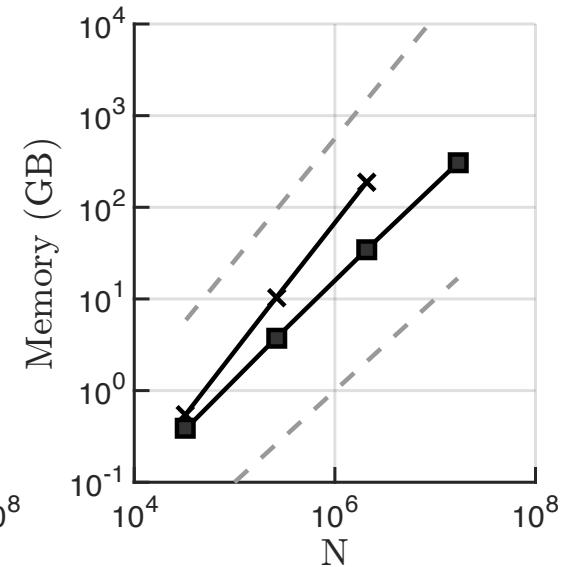
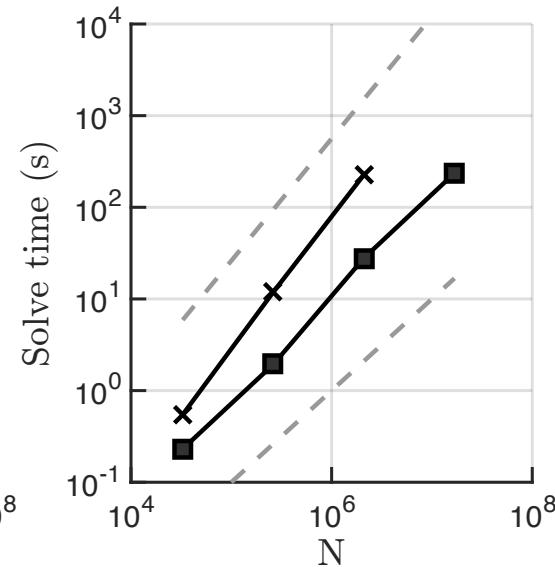
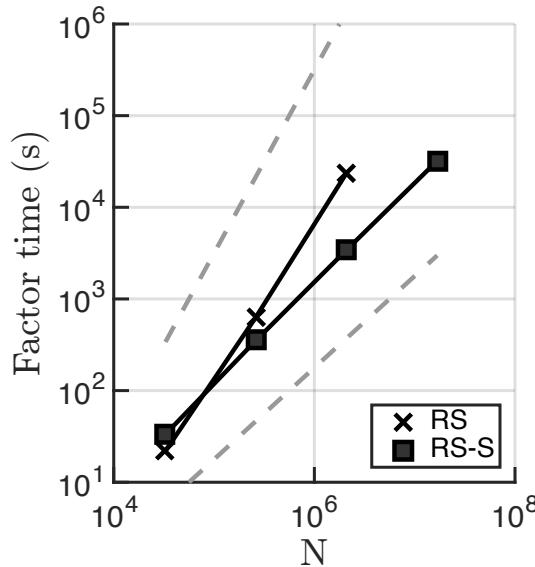
- Intel(R) Xeon(R) CPU E7-8890 @ 2.50GHz

$$\int_{\Omega} \frac{1}{4\pi|x-y|} u(y) dy = f(x), \quad x \in \Omega = [0,1]^3$$

- Methods
  - › Recursive skeletonization (RS)
  - › Strong recursive skeletonization (RS-S)

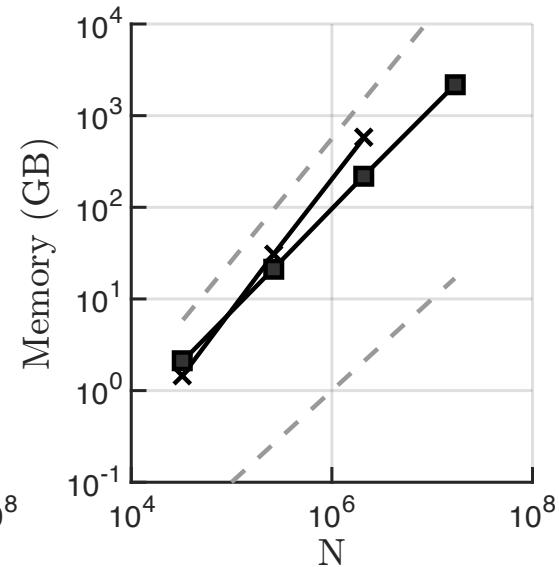
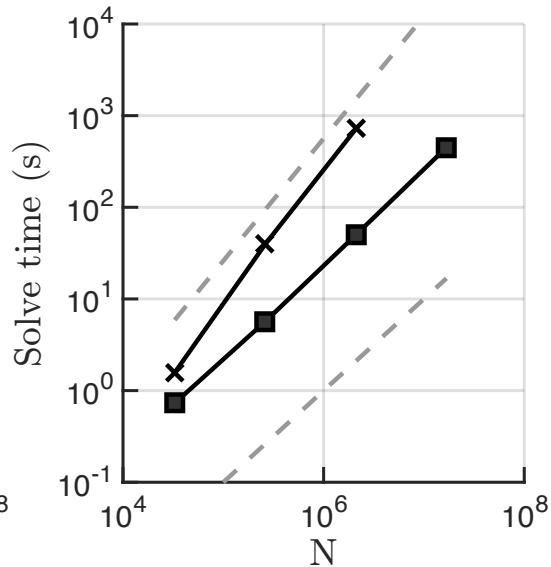
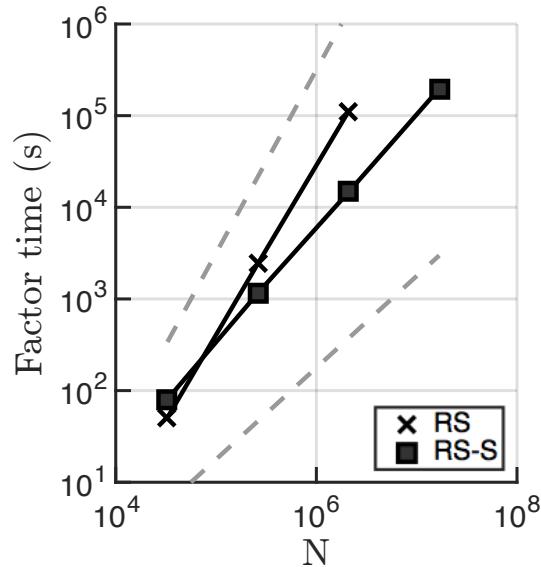
# Accuracy parameter $10^{-3}$

- Conjugate gradient converges to a relative residual norm of  $10^{-12}$  in about **10 preconditioned iterations**.



# Accuracy parameter $10^{-6}$

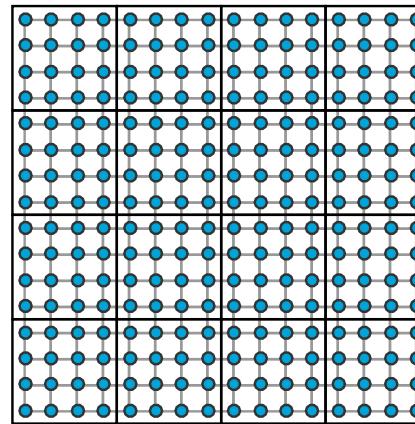
- Conjugate gradient converges to a relative residual norm of  $10^{-12}$  in about **3 preconditioned iterations**.



## Comments on results

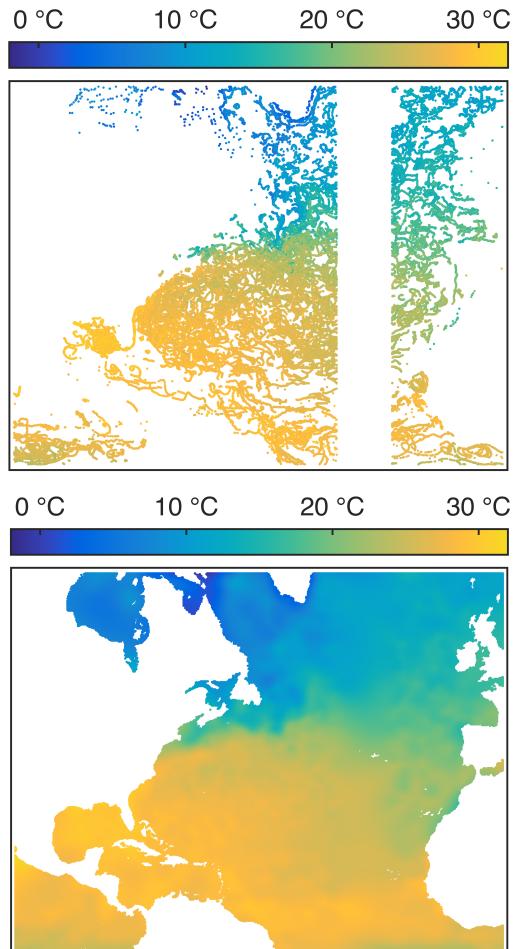
- Basic take-away
  - › New factorization can be viewed as an “inverse” form of the FMM based on recursive skeletonization
  - › Obtain an approximate multiplicative factorization of the inverse operator
  - › Exhibits **linear** scaling at preconditioner accuracy levels, avoiding excess rank growth from compressing near-field
  - › **Simple** to understand and to **implement**

# Conclusions



## Other work exploiting structure

- Updating hierarchical factorizations after local perturbations [Minden et al., 2016]
- Green's functions for time-stepping Maxwell's equations [Lo, Minden, & Colella, 2016]
- Maximum likelihood estimation for spatial Gaussian processes [Minden et al., 2016]
- Sparse canonical correlation analysis (ongoing, with Xiaotong Suo)



# Papers referenced in this talk

- Graph clustering
  - › References refer to bibliography section of  
Damle, Minden, & Ying, “Robust and efficient multi-way  
spectral clustering”, [arXiv:1609.08251](#)
- Hierarchical factorizations
  - › References refer to bibliography section of  
Minden, Ho, Damle, & Ying, “A recursive skeletonization  
factorization based on strong admissibility”, [arXiv:1609.08130](#)

# Acknowledgments

- V.M. is supported by a Stanford Graduate Fellowship and was previously supported by a U.S. Department of Energy Computational Science Graduate Fellowship under grant number DE-FG02-97ER25308.



Thanks for listening!